

Inertial Navigation using Atom Interferometry

Jimmy Stammers
Imperial College London

A dissertation submitted for ...

14476 (errors:1)

Abstract

This thesis describes work I did during my PhD...

Declaration

This dissertation is the result of my own work...

Jimmy Stammers

Acknowledgements

Some people worth thanking...

Preface

This thesis describes my research on various aspects of...

Contents

List of figures	xv
List of tables	xix
1 Introduction	1
2 Theory	3
2.1 Overview	4
2.2 Light-Matter Interactions	4
2.3 Laser Cooling of Rubidium-87	4
2.4 Raman Transitions in Rubidium-87	4
2.5 Light Pulse Atom Interferometry	4
2.5.1 The Double Interferometer	4
3 MOTMaster	5
3.1 Chapter Overview	5
3.2 Motivation	6
3.3 Interfacing wth Hardware	7
3.3.1 Hardware Abstraction	8
3.3.2 Voltage Pattern Generation	8
3.3.3 Timed Serial Communication	10
3.3.4 Voltage Acquisition	11
3.4 MOTMaster Sequences	11
3.4.1 Sequence Structure	12
3.4.2 Running a Sequence	13
3.5 Experiment Control Hardware	15
3.6 Experimental Sequence Overview	16
4 Cooling and Trapping in a MOT	19
4.1 Chapter Outline	19

4.2	The Navigator Vacuum Chamber	19
4.2.1	The 2D MOT system	22
4.2.2	The 3D MOT system	25
4.2.3	CCD Imaging	29
4.3	Generating MOT light	32
4.3.1	Absolute Frequency Reference	32
4.3.2	Cooling and Repump Light	35
4.3.3	Real-Time Control	36
4.4	Controlling the MOTs	37
4.4.1	Optical Fibre Network	37
4.4.2	Magnetic Field Control	38
4.5	Characterising the MOTs	40
4.5.1	3D MOT Loading Rate	42
4.6	Conclusion	43
5	Preparing Atoms for Interferometry	47
5.1	Chapter Overview	47
5.2	Cooling in Optical Molasses	48
5.2.1	Motivation for Launching	48
5.2.2	Frequency and Power Control	49
5.2.3	Launching in a Moving Molasses	49
5.2.4	Imaging the Atom Cloud over Time	52
5.3	State Preparation	56
5.3.1	Schemes for Preparation	57
5.3.2	Optically Pumping the Atoms	59
5.3.3	Including Microwave Transitions	61
5.3.4	Blow-Away	65
5.3.5	Conclusion	65
6	Acceleration-Sensitive Interference	67
6.1	Chapter Outline	67
6.2	Raman Optical System	68
6.2.1	Fringe Contrast Dependence	69
6.2.2	Raman Beam Collimator	72
6.2.3	Retro-reflection Assembly	76
6.2.4	The MEMS Accelerometer	81

6.3	Driving Raman Transitions	81
6.3.1	Frequency and Phase Control	82
6.4	Atom Detection	82
6.4.1	Optical System	82
6.4.2	Measuring the Interferometer Phase	82
6.5	Individual Pulse Characterisation	82
6.5.1	Velocity-Selective Pulse	82
6.5.2	Interferometer Pulses	82
6.6	Three-Pulse Atom Interference	82
6.6.1	Fringe Contrast vs. Launch Trajectory	82
6.7	Measuring Accelerations	82
6.7.1	Vibration Sensitivity	82
7	Outlook	83
7.1	Combining with classical accelerometers	83
7.2	Extending to sensitivity along three axes	83

List of figures

3.1	Schematic diagram of MOTMaster and hardware interface	7
3.2	Scripted pattern generation for an NI-HSDIO card	10
3.3	Timing diagram for serial communication	11
3.4	MOTMaster user interface	14
3.5	Control hardware schematic diagram	17
4.1	Magneto-optical Trap (MOT) system component diagram	21
4.2	Schematic for the 2D MOT	23
4.3	Optical components for the 2D MOT	24
4.4	Simulated field and field gradients for the 2D MOT quadrupole coils .	26
4.5	Measured magnetic field and field gradient for the 3D MOT coils . .	28
4.6	Simulated magnetic field for the 3D MOT quadrupole coils	28
4.7	Optical setup for CCD imaging	30
4.8	Integrated pixel counts as a function of CCD exposure time.	31
4.9	μ Quans Laser System Diagram	33
4.10	Saturated absorption spectroscopy of the Muquans master laser.	34
4.11	Error Signal for the μ Quans master servo.	35
4.12	Error Signal for the μ Quans Cooling laser.	35
4.13	Network of optical fibre splitters and aoms!s (aoms!ls) for MOT! (MOT!) light distribution	39

4.14 3D MOT atom number vs dispenser current with and without a push beam.	44
5.1 Molasses stage timing diagram	50
5.2 Beam configuration for a moving molasses	51
5.3 Atom cloud position after launching in a moving molasses	52
5.4 Temperature measurement using ballistic expansion	54
5.5 Atom cloud centre-of-mass over time	55
5.6 State preparation pulse sequence	58
5.7 State selection timing schematic	59
5.8 Ground state distribution during $ F = 2\rangle \rightarrow F' = 2\rangle$ pumping	60
5.9 m_F populations before and after $ F = 1\rangle \rightarrow F' = 0\rangle$ pumping	61
5.10 $ 1, m_F\rangle$ populations for increasing $ F = 1\rangle \rightarrow F' = 0\rangle$ pumping time.	62
5.11 Setup for Microwaves	63
5.12 Microwave Rabi oscillation between $ 1, 0\rangle$ and $ 2, 0\rangle$	64
5.13 Microwave transition spectrum	64
6.1 Simulated fringe contrast vs beam waist size	70
6.2 Expected contrast as a function of random phase contributions. This assumes that the phase imprinted on an atom during each interferometer pulse has an additional random component that is Gaussian distributed around 0 with a standard deviation of σ_ϕ . This random phase is also uncorrelated between each pulse so that the total can be obtained using Gaussian propagation of error. The dashed lines indicate the contrast for phase noise expected from conventional optics, which are usually engineered to a surface flatness of $\lambda/20$	71
6.3 Drawings of the components used in the Raman optics assemblies	74
6.4 Simulated wavefront distortion for longitudinal and transverse fibre misalignment	75

- 6.5 Thickness of the first **QWP!** (**QWP!**), measured by a white light interferometer. The value is given in nm as a difference from the mean thickness. The standard deviation of this thickness is 4.62 nm and a peak-to-valley (PV) of need number here. Equivalent surface data for the other **QWP!** and mirror were not provided by Light Machinery, but had a PV thickness variation of 19 nm and 9 nm respectively. 77
- 6.6 Automatic mirror alignment using the Nelder-Mead simplex algorithm. 80

List of tables

4.1	Table of parameters for each 3D MOT bias coil.	29
4.2	Typical optical and magnetic parameters for the Magneto-optical Traps (MOTs).	41

Chapter 1

Introduction

- atom interferometry experiments for precision measurements of inertial forces
- inertial navigation suffers from long-term bias drift
- recent experiments have demonstrated measuring acceleration in environments of interest to navigation

Chapter 2

Theory

- Describe general principles of light-matter interaction
- Specific cases for laser cooling (doppler/sub-doppler) and Raman transitions
- Lead into atom interferometry
- Perhaps split this into two shorter chapters

2.1 Overview

2.2 Light-Matter Interactions

2.3 Laser Cooling of Rubidium-87

2.4 Raman Transitions in Rubidium-87

2.5 Light Pulse Atom Interferometry

2.5.1 The Double Interferometer

Chapter 3

MOTMaster

3.1 Chapter Overview

The aim of this chapter is to provide a description of the MOTMaster software, which was developed from a pre-existing version during my PhD. The design of MOTMaster assumes very little about the particular experiment it is being used for, so much of the discussion in this chapter will be kept general. This chapter begins with a motivating the need to extend MOTMaster by developing a graphical interface to simplify the creation of experimental sequences, as well as implementing new methods of controlling hardware. This is followed by a description of how input and output channels are controlled using MOTMaster in Section 3.3. The structure of a MOTMaster sequence, along with how it runs an experiment is then presented in Section 3.4. Finally, the specific hardware used in this experiment and an overview of each major step of the experiment is given in Section 3.5.

3.2 Motivation

In the initial stages of my PhD, I decided to use Cicero Word Generator [3] to control the hardware for the experiment. This is a graphical-based control system developed by Wolfgang Ketterle's group at MIT, which was designed for controlling atomic physics experiments using National Instruments hardware. Over time, as the experiment became more complex, it started to become apparent that Cicero was not suited to meet all of our requirements for control software. This was most evident in the control of the M-Squared Raman laser system. Cicero also takes an appreciable amount of time (around 300 ms) to re-calculate the experiment sequence between each shot. Since the design of Cicero was aimed at controlling experiments that take many seconds per cycle, this dead time between each cycle is not significant on those time scales. In contrast, each cycle of this experiment takes around 250 ms. This unnecessary dead time needed to be addressed if we hoped to improve the repetition rate.

After it became clear that a potentially large amount of work would be needed to improve Cicero, I decided that it was worth moving to a new control system. A collection of programs, named EDMSuite, has been developed by people in Centre for Cold Matter (CCM) to control a range of experiments within the group. One application, MOTMaster, was designed to control and acquire data from experiments investigating cold atoms trapped in a MOT. However, its method of structuring experimental sequences was inconvenient, as it lacked an intuitive graphical user interface. During the process of switching to using MOTMaster to control the experiment, I designed a graphical method of structuring sequences, which functioned identically on a device level to the original method of defining sequences. In addition to this, I included an interface to the M Squared laser system, so that it could be controlled using MOTMaster. A schematic of the structure of MOTMaster and how it interfaces with hardware

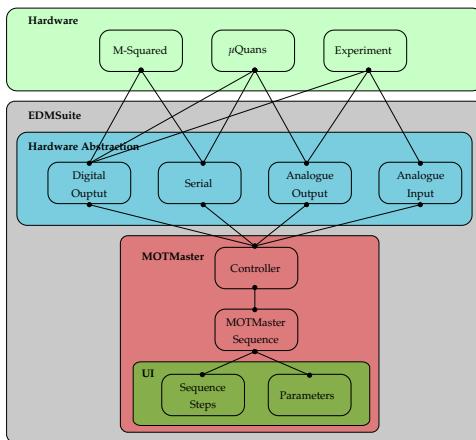


Figure 3.1: Schematic diagram of MOTMaster and the hardware it controls. A sequence is built using the user interface, which then uses separate modules to communicate to the hardware. In this way, an experiment can be controlled without requiring specific knowledge of the hardware.

is shown in Figure ???. It is designed so that experiments can be controlled without requiring specific details about the hardware in use.

3.3 Interfacing wth Hardware

The majority of the experimental hardware is controlled using analogue and digital voltages that are generated by Data Acquisition (**DAQ**) cards manufactured by National Instruments. MOTMaster is compatible with cards that use either the NI-DAQmx or NI-HSDIO device drivers. These are used to configure the generation or acquisition of digital or analogue voltage waveforms. By design, they are capable of precisely timing and synchronising their I/O across multiple devices. Most components in the experiment rely on this precise timing to function correctly. Other devices, where timing accuracy is less critical, are controlled by sending or receiving data using serial communication. This has the advantage of allowing more structured command beyond analogue or digital voltages, but the communication speed of the serial channel limits the accuracy of the execution time.

The following section describes the low-level interface between MOTMaster and the experimental hardware. It begins by introducing the concept of hardware abstraction in Section 3.3.1. This is followed by a more detailed discussion of how each type of control is implemented. Section 3.3.2 describes how analogue and digital output waveforms are generated. Section 3.3.3 outlines serial communication, along with a method for triggering this communication during an experiment. Finally, this section concludes with a discussion on acquiring analogue input data, which is given in Section 3.3.4.

3.3.1 Hardware Abstraction

When designing software, it is often useful to structure a program in such a way that modules which make use of other components do not need to know about their specific implementation in order to use them. This approach means that the submodule can be modified without harming the compatibility of these two components. In the context of experimental hardware, this is equivalent to requiring that changing specific components, for example the Voltage-Controlled Oscillator (**VCO**) that generates the RF power for an Acousto-optic Modulator (**AOM**), will not stop the experiment from working. This is done using abstract representations of the hardware, in the form of input and output channels that are used to communicate to each device.

3.3.2 Voltage Pattern Generation

Analogue Outputs

All the analogue outputs controlled using MOTMaster are done using the NI-DAQmx software. Each output uses a Digital-to-Analogue Converter (**DAC**) to convert a

floating-point number into an analogue voltage. To generate a sequence of voltages across multiple channels, the NI-DAQmx driver allocates a block of memory on the **DAQ** card for each output channel. This memory acts as a first-in first-out (FIFO) buffer for data streamed to it from a computer. The output of each channel is synchronised to a clock signal, so that every time a rising edge occurs on the clock, the voltage at each output transitions to the value corresponding to the next value in its corresponding buffer. Channels across multiple **DAQ** cards can be synchronised by sharing a clock signal, which can be done using the bus that connects cards in a PXI-e chassis. Additional cards can also be configured to trigger the start of their output at the moment they receive the first clock pulse, rather than waiting for a software trigger from the computer.

Digital Outputs

Digital outputs from NI-DAQmx cards are generated in much the same way as analogue voltages, except for the fact that they only take two values corresponding to either a low (0 V) or high (3.3/5 V) level. Additionally, **DAQ** cards which use the NI-HSDIO driver can be used. These cards can be sampled at much higher rates than NI-DAQmx ones. For instance, the NI-HSDIO PXI-6541 card can generate digital voltages at sample rates up to 50 MHz. Rather than writing the pattern as an array of values at each clock cycle, the sequence is segmented into smaller patterns during which the state of each channel is constant, as illustrated in Figure 3.2. NI-HSDIO cards can be scripted to generate each of these patterns for the appropriate number of clock cycles.

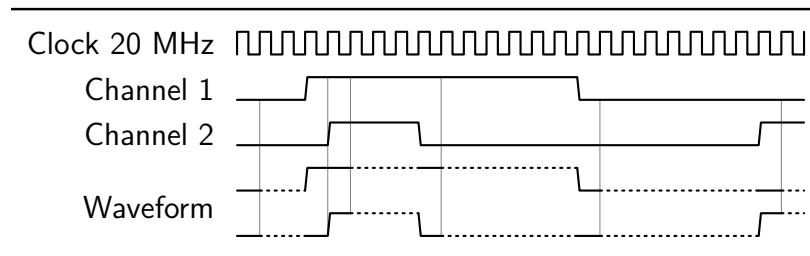


Figure 3.2: Scripted pattern generation for an NI-HSDIO digital output card. A pattern is split into segments which correspond to a duration for which all the channels output a constant value. Each of these smaller waveforms are written to the on-board memory, along with a script that instructs the card to output each pattern for the required number of times to reconstruct the original sequence. By reducing the amount of memory required to define the sequence, a faster clock frequency and hence timing resolution can be used to output digital control signals.

3.3.3 Timed Serial Communication

Serial communication is used to control devices which require more complex control than is possible using analogue or digital voltages. This increase in complexity comes at the cost of slower response times, because it takes longer to communicate an array of bytes than to change the voltage across an output terminal. Using the NI-VISA driver, the output of serial data can only be timed using a software clock on a computer, which is more prone to jitter than a hardware clock. One way to improve the synchronisation between serial data and hardware timed outputs is to use extra hardware to trigger the transmission of serial data. If the trigger is timed using the same clock as other outputs and the transmission delay is accounted for, then serial data can be output more synchronously. The scheme for timing serial messages is shown in Figure 3.3. Serial messages are stored as strings on the computer and a counter channel is configured so that every time it detects a rising edge, the computer outputs the next message. This counter is connected to a digital output channel, so that it acts as a trigger for the serial data output. Using this method, multiple serial messages can be sent to one device during a sequence even for devices which have no means of storing commands.

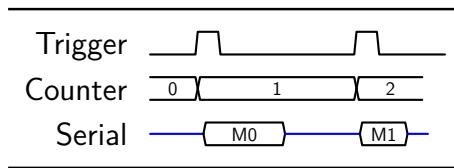


Figure 3.3: Timing diagram for serial communication. A counter channel is configured to count edges from a digital output channel. Every time it sees a rising edge, it triggers the output of the next message on each serial channel from the computer. Multiple messages can be communicated during a single sequence without the need for software timing.

3.3.4 Voltage Acquisition

Analogue input channels are configured in a similar way to analogue output channels. A block of memory is allocated on the DAQ card for each input channel. Once the card is triggered to start acquiring, an Analogue-to-Digital Converter (ADC) converts the voltage across the input into a digital value at every rising edge of the clock signal. Once the sequence has finished, or the buffer has been filled, the card streams this data to the computer.

3.4 MOTMaster Sequences

In addition to interfacing with control hardware, MOTMaster is used to define the structure of experimental sequences. In earlier versions of MOTMaster, sequences were defined using functions within a C# source file. To run an experiment, MOTMaster compiled this file to build the voltage patterns and wrote them to the hardware. Whilst this had little overhead in resources needed to build and run a sequence, modifying and debugging sequences was much more time consuming. Taking inspiration from Cicero, the user interface of MOTMaster was redesigned so that sequences could be expressed graphically. They are then built using the same functions as before, so that from the point of view of the hardware, the two methods of control are equivalent.

3.4.1 Sequence Structure

A MOTMaster sequence is composed of a list of sequence steps, which define the state of the control hardware over a discrete amount of time. Each step contains the following properties:

- Name: A descriptive name for the step.
- Duration: duration of the step, which must be an integer multiple of the timebase (e.g. 10 µs for a 100 kHz sample clock frequency).
- Serial Channel: A serial message encoded as a string of text.
- Digital Output Channel: High (3.3 V) or Low (0 V)
- Analogue Output Channel: Single value, step or ramp the output from a start to end value, or output an arithmetic function over time.

where each individual output channel has its own property.

A sequence step is useful to represent a single action, so that each stage of the experiment, for example the initial MOT loading phase, is composed of multiple steps. Numerical values, such as analogue voltages or times, can be represented by named parameters. The value of a parameter can be updated between each cycle of the experiment, so that MOTMaster can implement a scan by iterating a parameter through a range of values. The sequence steps are also used to define when to acquire from the analogue inputs. A specific digital channel, named acquisitionTrigger, is reserved as a start trigger for the acquisition. This channel is also used to define the length of time over which to acquire data. Analogue data acquisition is triggered at the start of the step where this channel goes high and stops when it goes low.

3.4.2 Running a Sequence

MOTMaster is designed to run in two modes, referred to as repeat and scan. The distinction between these is that the repeat mode does not need to recreate a sequence between each cycle. Before MOTMaster starts controlling the experiment, the sequence is built once and the output hardware is configured to regenerate their patterns. This reduces the delay between each cycle, which is largely a result of the time needed to process acquired data and reconfigure the control hardware. In contrast, scan mode varies a parameter during each cycle, so additional time is required to rebuild the sequence and write to each **DAQ** card. Aside from this, these modes operate equivalently.

At the start of an experiment cycle, the hardware is initialised and timing properties, such as the trigger and sample clock for each **DAQ** card is set. An example of a sequence as represented in the user interface is shown in Figure 3.4. This is converted into the analogue and digital voltage patterns for each **DAQ** card. The required buffer for the analogue input data is calculated based on the state of the acquisitionTrigger channel. If any serial commands are used, the timing properties of the counter channel are configured, similarly to the rest of the **DAQ!** (**DAQ!**) hardware. The sequence is started by sending a software trigger to one output card, which is configured to export its start trigger to the other cards. This ensures that start of the output of each card is synchronised.

After the sequence has finished, any acquired data from the analogue input channels is streamed to the computer. The data per channel are segmented into arrays that were acquired during each sequence step, before additional post-processing if required. Finally, the hardware is reset to its initial state, before starting the next experiment cycle.

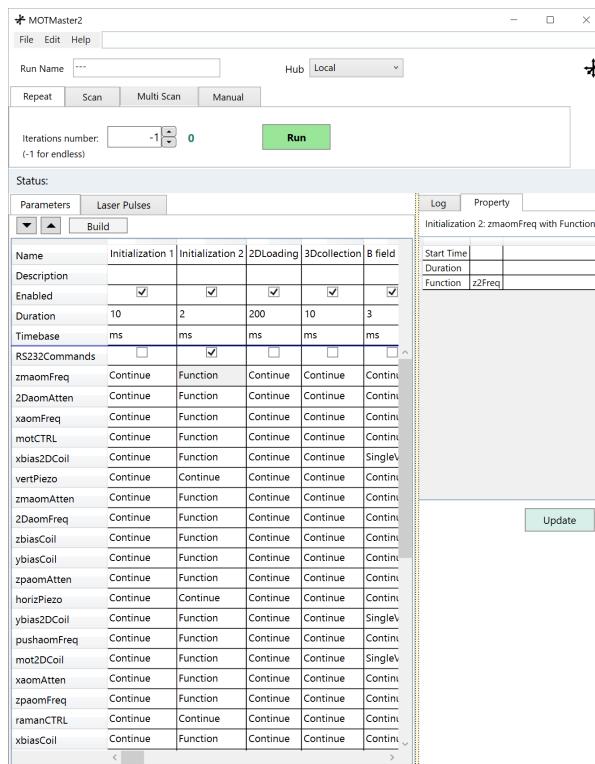


Figure 3.4: A sequence as represented in the MOTMaster user interface. Each step defines a duration and properties for each output channel. The sequence can either be run repeatedly, or configured to iterate through values of a chosen parameter.

3.5 Experiment Control Hardware

In the preceding sections, the discussion of MOTMaster has been presented without referring to specific hardware used in this experiment. Subsequent chapters will introduce components of the experiment that are controlled by a computer, but it is worth introducing the hardware used to implement this control. A diagram of the control hardware is shown in Figure 3.5. All of the **DAQ** cards are housed on a PXIe-1073 chassis, so that timing signals such as start triggers and sample clocks can be shared on the PXI backplane. The analogue output signals are generated on a PXI-6723 card. This contains 32 analogue output channels and the output of each is generated using a 13 bit **DAC**. Over the maximum voltage range of $\pm 10\text{ V}$, this corresponds to an output quantisation of 2.44 mV , which did not limit the precision of any analogue control in the experiment. The analogue output pattern is sampled at a frequency of 100 kHz , which gives a minimum resolution of $10\text{ }\mu\text{s}$. Any jitter on this sample clock did not produce any noticeable effects during the experiment.

Two cards on the chassis are able to acquire data from analogue inputs. The first is a PXIe-6341, which has 16 input channels, each with a 16-bit **ADC**. In addition to this, a counter channel on this card was used to trigger the output of serial messages. During the preliminary stages of the experiment, this bit-depth was sufficiently large to prevent quantisation effects becoming significant. However the AI-Q-2010 MEMS accelerometer used in the experiment, discussed further in Section 6.2.4, has an equivalent voltage noise below this quantisation level. Therefore, a PXI-4462 card, which contains 4 24-bit analogue input channels, was added. This card is used to acquire data from devices where the higher voltage resolution is desirable — namely, the MEMS accelerometer and the photodiode used to detect the population of atoms in each state after interference.

Digital output signals are generated using a PXI-6541 card. Unlike the others, this card is controlled using the NI-HSDIO driver. With a maximum sampling frequency of 50 MHz, this card is capable of generating digital signals at a much higher rate than the PXIE-6341, which also contains digital output channels. However, the PXIE-6341 card only contains 8 digital channels that can be timed using a hardware clock, fewer than required to control the entire experiment.

Two components of the experiment are controlled during the experiment using serial communication. The first of these is an interface to the Direct Digital Synthesiser (DDS) on the μ Quans laser which control the frequency of the cooling and repump lasers and is controlled in real-time during the experiment. This communication protocol is described in further detail in Section 4.3.3. Finally, MOTMaster is configured to remotely connect to the M Squared laser, so that it can control all the parameters necessary to drive Raman transitions during the experiment. This is done by sending structured JSON messages that contain commands to implement this control. More detail on how this is used in the experiment is given in Section ??.

3.6 Experimental Sequence Overview

The experiment can be broken down into the following stages:

- Loading: Atoms are loaded from the 2D MOT into the 3D MOT.
- Molasses: Atoms are released from the trap and cooled further in an optical molasses.
- State Preparation: A sequence of optical and microwave pulses are used to prepare atoms with a narrow velocity spread in the $|1, 0\rangle$ state.

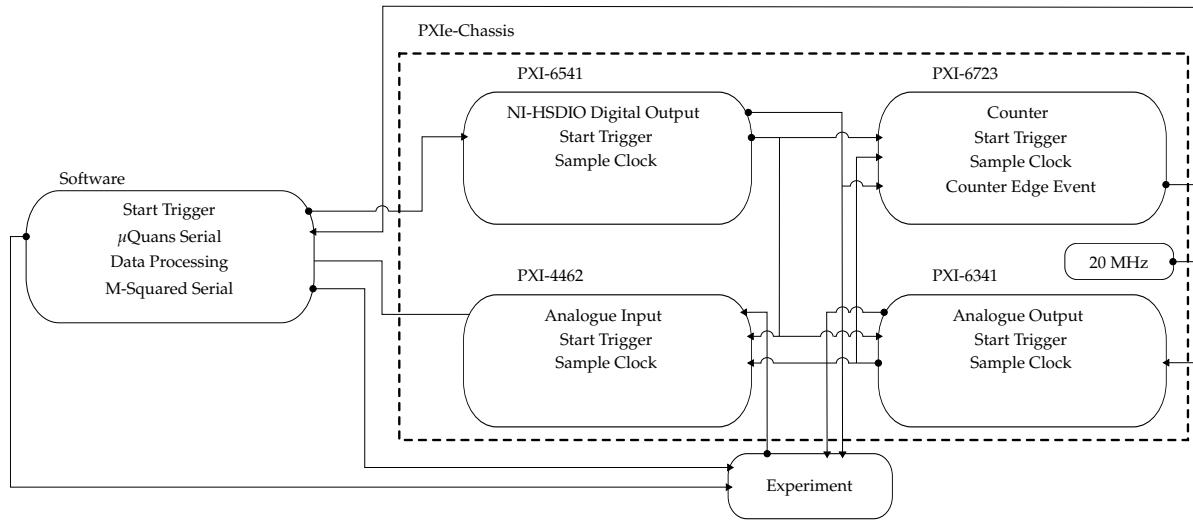


Figure 3.5: Schematic diagram of the control hardware. The PXle chassis contains the DAQ cards which generate the analogue and digital waveforms used to control other devices. Signals are routed between the cards to synchronise their operation. Serial communication to the μ Quans laser (Section 4.3), the M-Squared laser (Section 6.3.1) and the CCD camera (Section 4.2.3) are used for their control.

- Interferometry: A $\pi/2 - \pi - \pi/2$ sequence of laser pulses drive Raman transitions between the atoms.
- Detection: Two laser pulses are used to measure the number of atoms in $|F = 2\rangle$ and the total number, respectively. From these measurements, the interferometer phase difference can be inferred.

Chapter 4

Cooling and Trapping in a MOT

4.1 Chapter Outline

This chapter presents a description of the components of the experiment which are used to trap and cool atoms in a MOT. An outline of the hardware used to create both the 2D and 3D MOTs is presented in Section 4.2. Following this is a description of the μ Quans laser system, in Section 4.3, which generates the light used to cool and trap atoms. The hardware used to control the frequency and power of each MOT beam, as well as the required magnetic fields, is given in Section 4.4. Finally, this chapter concludes with a presentation of a measurement of the 3D MOT loading rate, which serves as a figure-of-merit for the performance of the MOT system.

4.2 The Navigator Vacuum Chamber

The vacuum chamber, along with the components mounted to it, make up the majority of the hardware used in the preliminary trapping and cooling stages of the experiment.

The chamber is made of 316L stainless steel which helps to minimise the influence of stray magnetic fields that could otherwise affect the performance of the MOT. The chamber contains 16 DN40 ConFlat ports arranged on the edges of three octagons, one in each cartesian coordinate plane. Six of these ports are used to provide optical access for the 3D MOT, and a further four are used to connect to the 2D MOT system, a gate valve, the NexTorr pump and to provide power to the in-vacuum 3D MOT coils. Two of the remaining ports are used to provide optical access for the imaging systems — a CCD camera and a photodiode — or a microwave horn, depending on the specific experiment. The last four ports are not used for additional optical access since their line of sight to the atoms is obscured by the MOT coils. Two DN63 ports lie along one axis, which is conventionally taken to be the x axis. These ports are used to mount the optics for driving Raman transitions and as such, defines the axis along which the atom interferometer is sensitive to accelerations¹.

A diagram of the vacuum chamber and the main MOT components is shown in Figure 4.1. The chamber is pumped down to a pressure of around 5×10^{-10} mbar using a NexTorr D100-5 pump. This is a composite system consisting of neg! (neg!) and an ion pump. The neg! is a porous sintered zirconium (St 172) element, which reacts with chemicals such as hydrogen, water, nitrogen, oxygen and hydrocarbons. Most of these were removed during the initial baking and roughing pump stages. Under uhv! (uhv!) conditions, the largest contributor to the pressure is hydrogen which the neg! can pump at a speed of 100 l s^{-1} . Any species that are not absorbed by the neg!, in particular Rubidium, are pumped by the 5 l s^{-1} ion pump.

¹For more information about the Raman optical system, refer to Section 6.2

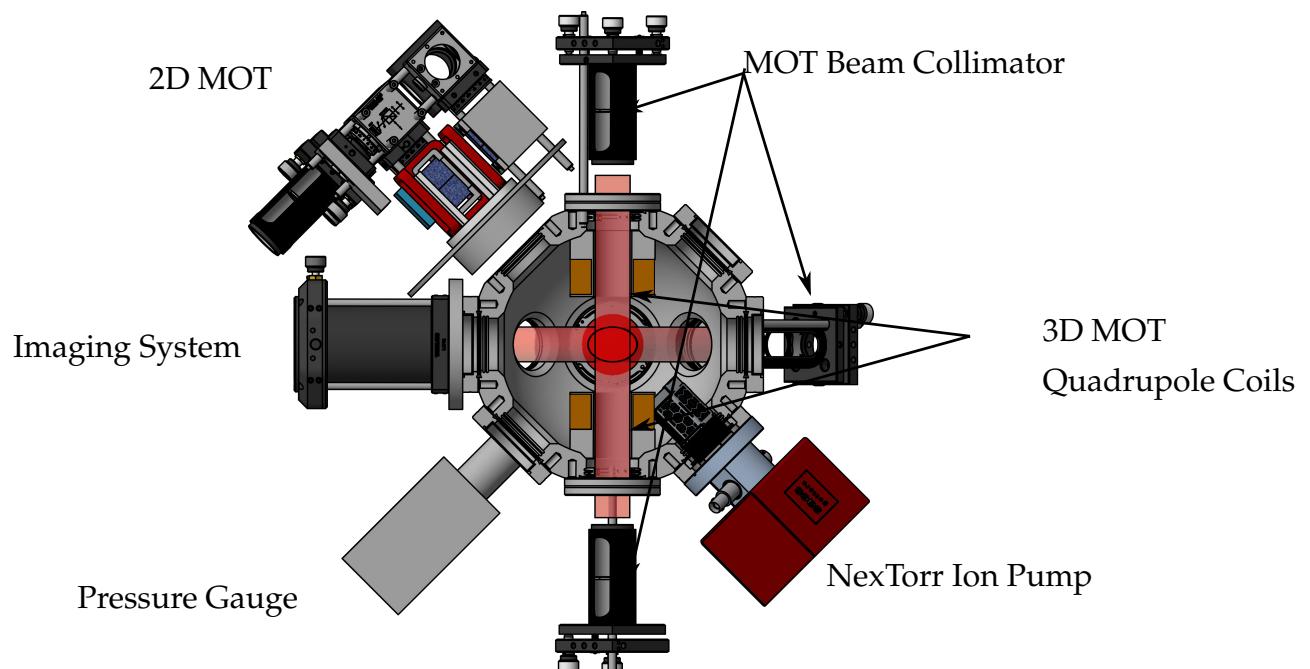


Figure 4.1: A diagram of the main components on the vacuum chamber used for the MOT systems. Rubidium atoms are dispensed and loaded into the 2D MOT before being pushed into the main chamber and collected in the 3D MOT. A set of 6 beam collimators provide the light necessary to slow and cool atoms, which are trapped using the spherical quadrupole field generated by the illustrated coils. Not shown are additional bias coils along each MOT beam axis to null stray fields at the centre of the chamber.

4.2.1 The 2D MOT system

Initially, atoms were loaded into the 3D MOT from a background vapour. Whilst this was a relatively simple scheme, a very high partial pressure of Rubidium is required to achieve a fast loading rate (see Section 4.5.1). However, this was undesirable from the point of view of the atom interferometer, as it would result in a poor fringe contrast due to collisions with background atoms and a worse signal to noise ratio, i.e. the ratio of the number of atoms in the interferometer to the total number of detected atoms. Replacing the source of atoms for the 3D MOT with a side-arm to function as a 2D MOT [2] satisfied the two requirements of a fast loading rate and low base pressure in the main chamber.

A diagram of the light and magnetic fields required to produce a 2D MOT is presented in Figure ???. It is similar to the 3D MOT, with the main exception being that only 4 beams are used to cool the atoms along 2 orthogonal axes. In addition to this, its design is focused towards the production of a large flux of cold atoms which can be subsequently loaded into a 3D MOT. For instance, the beams used to cool the atoms are collimated to a large waist size and the coils are designed to give a cylindrical quadrupole field with a line of zero magnetic field along the axis of symmetry. Along this axis, the atoms are free to move which results in an atomic beam. To improve the collimation of this atomic beam, a larger radial field gradient than usually used in a 3D MOT system is used to increase the radial confinement of atoms. In addition, a pinhole is placed at the exit of the cell, so that atoms with a high radial velocity component will miss the aperture. This pinhole also greatly reduces the conductance between the 2D MOT cell and the main chamber, which means that a comparatively high background pressure (hence, loading rate) can be maintained in the 2D MOT cell, without greatly increasing the pressure in the main chamber. The pinhole is drilled into a silicon plate, which is used to partially reflect a beam that propagates along the

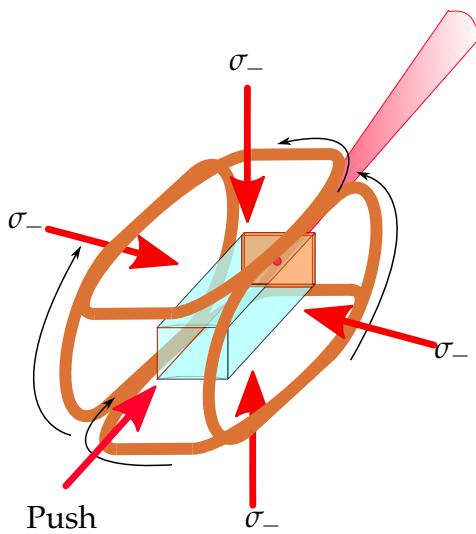


Figure 4.2: Schematic diagram for the 2D MOT. ^{87}Rb atoms are trapped and cooled along the 2 axes orthogonal to the long axis of the source cell. The arrows indicate the direction of the current through each coil, so that each circularly polarised beam drives σ_- transitions for an atom moving in the opposite direction. A linearly polarised push beam propagates along the remaining axis and is partially reflected by the silicon wafer at the opposite end. This provides a small amount of axial cooling and the imbalance of radiation pressure pushes atoms out of the cell. The pinhole at the other end prevents atoms with a high transverse velocity from leaving the cell.

central axis. This creates an unbalanced molasses that cools atoms with a large axial velocity and due to the differing radiation pressure, pushes atoms out through the pinhole. By slowing a larger proportion of atoms to within the capture velocity of the 3D MOT, this configuration, referred to as a 2D MOT, loads a 3D MOT faster than the 4-beam counterpart.

A schematic of the optical components for the 2D MOT is presented in Figure ???. The cooling light originates from a single fibre, which is collimated using two aspheric lenses to a beam waist of 9.5 mm and linearly polarised before being evenly split into two beams using a Half-wave Plate (HWP), one for each cooling axis. Each beam passes through a beam-splitter and a prism mirror, to increase the volume covered by the 2D MOT beams and is circularly polarised by a pair of Quarter-wave Plate (QWP) before entering the ar-coated glass cell. On the opposing side of the cell, a $25\text{ mm} \times 35\text{ mm}$ retro-reflecting mirror is used to provide the counter-propagating

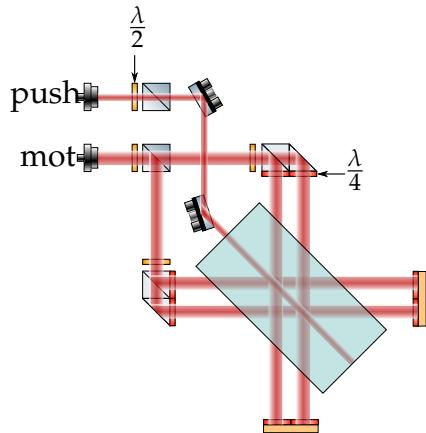


Figure 4.3: Optical components for the 2D MOT. The light for the MOT is split into two equal portions using a HWP and PBS. Along each axis, the beam passes through a PBS and a prism mirror to increase its spatial extent. The beam is circularly polarised before entering the cell and retro-reflected by a mirror coated with a QWP. The push beam is collimated from another fibre input and linearly polarised before entering the cell along the longitudinal axis.

MOT beam. This is coated with a layer of quartz to form a QWP, so that the reflected beam has the required polarisation for cooling atoms along that axis. The push beam is created using a second fibre input and a fixed collimator to give a beam waist of 1.5 mm. This is mounted onto a 1 in kinematic mount so that the alignment of the push beam with respect to the 0.7 mm pinhole at the other end of the cell can be optimised. A linear polariser is placed here to reduce the effect of polarisation drift on the axial cooling of the 2D MOT. The cell, manufactured by ColdQuanta, has dimensions of 30 mm × 30 mm × 44 mm and is specifically designed for creating a 2D MOT and contains two rubidium dispensers composed of rubidium chromate (RbCrO_4) and a reducing agent. These were activated by passing a large current through them to remove a thin oxidation layer. To produce rubidium, a current is passed through the dispenser to trigger an electro-chemical reduction reaction. The reducing agent also acts to sorb the unwanted products, so that they do not contaminate the cell and only rubidium evaporates from the dispenser.

The cylindrical quadrupole field is generated by a set of coils that are manufactured by ColdQuanta. Their geometry is such that the axis of zero magnetic field coincides with the central longitudinal axis of the source cell. These produce a radial field gradient per current of $20 \text{ G cm}^{-1} \text{ A}^{-1}$. A simulation of the field gradient along each axis, along with the magnetic field in each plane of symmetry is shown in Figure 4.4, which indicates that the field gradient is very uniform across the centre of the cell. In addition to this, each MOT axis has a separately controlled pair of coils in a Helmholtz configuration to provide a bias field that cancels nulls the field along the 2D MOT axis. These coils are made of 20 turns of wire that has a core diameter of 1 mm and produces a field per current of 1 G A^{-1} .

4.2.2 The 3D MOT system

The main chamber contains the apparatus that is used to make a 3D MOT. A diagram of the optical and magnetic fields used to create this MOT is presented in Figure ???. Each MOT beam is created using a collimator that collimates a beam of light from a Polarisation-Maintaining (PM) fibre. These use a lens with a nominal focal length of 75 mm to produce a collimated beam with a waist size of 7.5 mm. At the output of each collimator is a QWP, with its slow axis oriented at a 45° angle to the fast or slow axis of the fibre, depending on the particular collimator, to produce either left- or right-handed circularly polarised light. The MOT beams are oriented so that their intersection is at the centre of origin of a cartesian coordinate system. Since the MOT forms at the position where the magnetic field is zero, this system is set up so that beams overlap at the centre of the chamber, which is equidistant from each MOT coil. The MOT beams along the axial direction of the quadrupole field (conventionally referred to as the \vec{z} direction) are orthogonally polarised to the others along the \vec{x} and \vec{y} directions. This takes into account of the different direction of the magnetic field

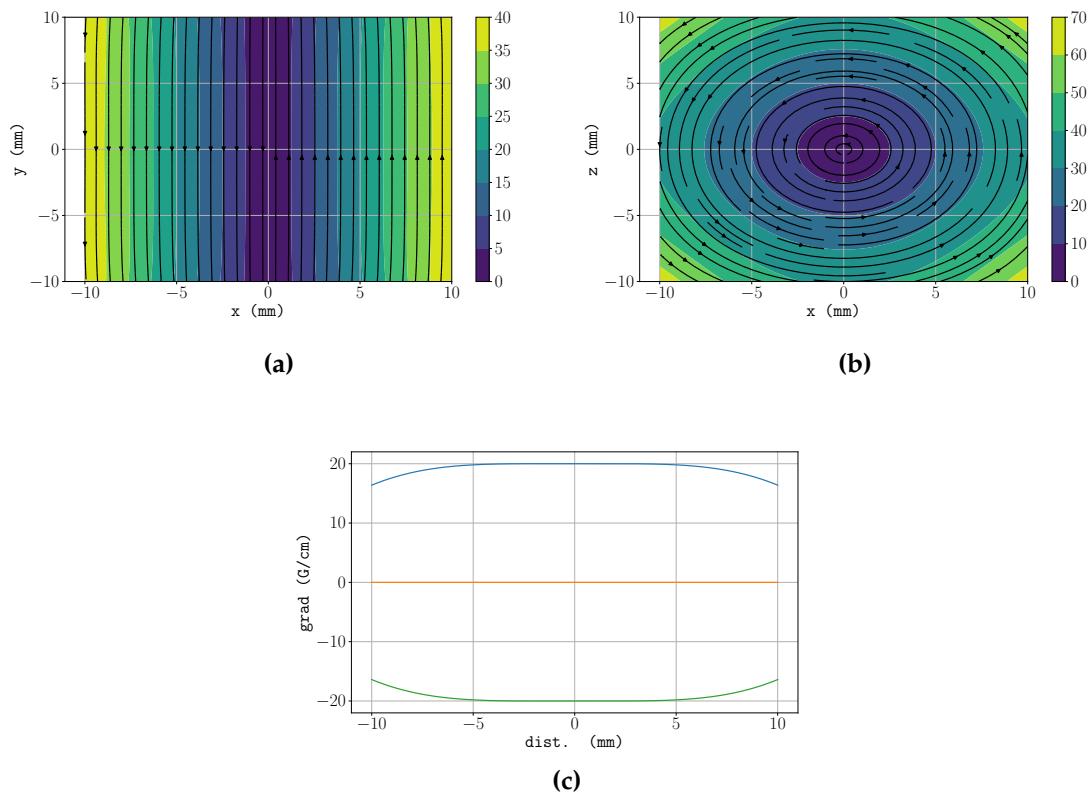


Figure 4.4: Simulated field and field gradients for the 2D MOT quadrupole coils. In this coordinate system, the 2D MOT cools and traps atoms in the \vec{x} and \vec{z} directions. The simulation was performed using a nominal current of 1 A, which corresponds to a current density in each coil of 7.78 A mm^{-2} . The magnitude of the magnetic field (in units of G) and its direction in the axial and radial planes of symmetry are shown in (a) and (b), respectively. (c) shows the field gradient components $\partial_x B_x$ (blue), $\partial_y B_y$ (orange) and $\partial_z B_z$ (green) along their corresponding axes.

gradient along each axis (see below), so that the polarisation of each MOT beam is such that an atom moving away from the centre of the trap is slowed and optically pumped into a state which feels a conservative potential.

Spherical quadrupole magnetic field

The magnetic field for the 3D MOT is created by a pair of coils in an anti-Helmholtz configuration. Each coil consists of rectangular wire coated in a 35 μm thick layer of Pyre-M.L, a UHV-compatible polyamide which provides a layer of insulation between each loop. The wire has a cross-section of dimensions 1.1 mm \times 1.1 mm, with 20 axial loops and 12 radial loops. The inner diameter of the coil is 25.4 mm, to allow for optical access of the \vec{z} -axis MOT beams, and the maximum diameter is 59.2mm – small enough that the coil could be inserted into the chamber through the DN63 CF ports. The coils are mounted to the chamber using groove grabbers which clamp into grooves inside the wall of the DN40 ports. The formers are designed to maximise the surface area between them and the chamber, which in turn aims to maximise the rate at which heat is dissipated from the coils. Once mounted, the distance between the innermost loops is 70 mm. Figure 4.5 shows the magnetic field, measured using a Hall probe, along the axis of symmetry for each of the coils with a current of 2.53 A. This is in close agreement with the value expected from simulating the field produced by the coils - the magnetic field in the axial and radial planes of symmetry and shown in Figure 4.6a and Figure 4.6b, respectively.

Bias Coils

Three orthogonally arranged pairs of Helmholtz coils are used during the experiment to null the magnetic field at the centre of the chamber. This is required for strong sub-Doppler cooling of the atom cloud in an optical molasses where the presence

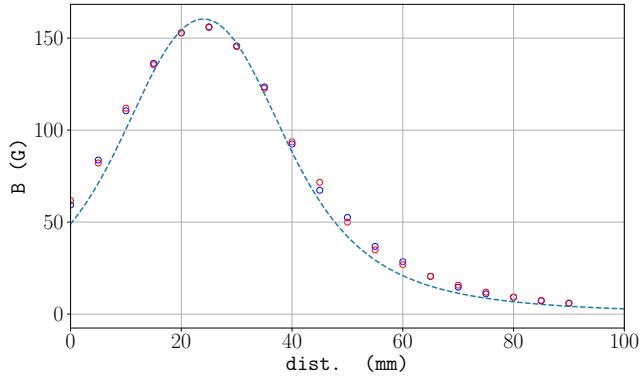


Figure 4.5: Measured magnetic field and field gradient for the 3D MOT coils. **(a)** shows the axial magnetic field for the two coils as measured using a Hall probe. The dashed line is the axial field as calculated from the simulation shown in Figure 4.6.

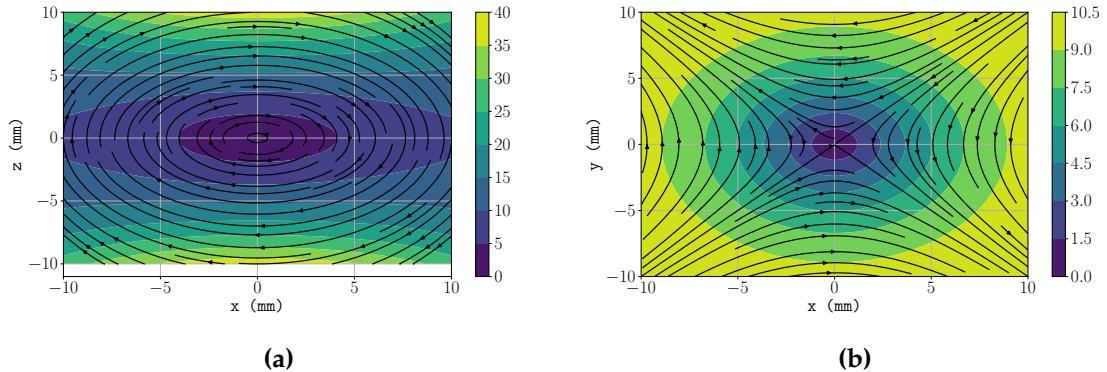


Figure 4.6: Simulated magnetic field for the 3D MOT quadrupole coils. In this coordinate system, the axial direction is defined as the \vec{z} axis. The simulation was performed using a nominal current of 2.53 A, which corresponds to a current density in each coil of 1.33 A mm^{-2} . The magnitude of the magnetic field (in units of G) and its direction in the axial and radial planes of symmetry are shown in **(a)** and **(b)**, respectively.

Axis	a (mm)	r_i (mm)	r_o (mm)
\vec{x}	88	105	115
\vec{y}	132	178	188
\vec{z}	116	123	133

Table 4.1: Table of parameters for each 3D **MOT** bias coil. a denotes the axial separation between each coil, r_i and r_o are the inner and outer radii. Each coil was wound to give 50 loops (5 axial and 10 radial turns).

of a magnetic field reduces the cooling efficiency [11]. These coils are also used in subsequent stages of the experiment to provide a bias field along the appropriate axes during state preparation, interferometry and state detection. Each coil was wound using 1 mm thick wire and consisted of 5 axial and 10 radial turns. For a pair of coils in Helmholtz configuration, the magnetic field gradient at the centre is minimised when the axial separation a is equal to the coil radius r , but the geometry of the vacuum chamber meant that it was not possible to satisfy this condition. The radii and axial separations of each coil pair is presented in Table 4.1.

4.2.3 CCD Imaging

The characterisation of the **MOT** stage of the experiment was achieved using a CCD camera to spatially resolve the atom cloud by detecting light emitted from the atoms during resonance fluorescence. Measuring the spatial distribution of atoms is useful, for example in estimating the temperature of the cloud by measuring the rate of thermal expansion. Figure 4.7 shows a diagram of the apparatus used for imaging. A pair of 125 mm and 50 mm focal length lenses are used to image the cloud onto a Pike F505-B CCD camera, which has a maximum resolution of 2452×2054 pixels. The pixel size in the object plane was measured by placing a ruler in that plane, giving a calibration factor of 5.1 pixel/mm. A bandpass filter is placed between the lens and

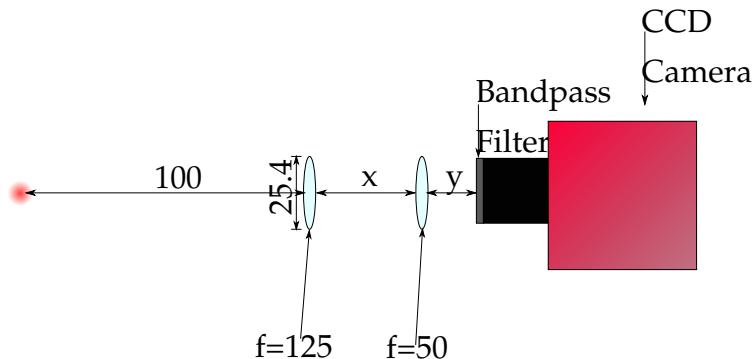


Figure 4.7: Optical setup for CCD imaging. Two lenses are used to magnify the image of the atom cloud on the CCD. A bandpass filter is placed in front of the sensor to block out background light at wavelengths other than 780 nm. All lengths are given in millimetres.

the CCD which transmits 780 nm light at an efficiency of 60% and blocks background light at other wavelengths.

As well as being used to gain information about the spatial distribution of atoms, the CCD camera can be used to make an estimate of the number of atoms N_a by measuring the amount of light emitted during resonance fluorescence. However, it is worth noting that this is not a reliable method for determining the number of atoms in a MOT. Typically, the density of atoms is such that the incident light is significantly absorbed by the outer atoms, resulting in a weaker intensity close to the centre. Consequently, the assumption of a constant scattering rate per atom is not valid and leads to an under-estimate of the atom number. A more accurate measurement can be performed by measuring the amount of light absorbed by the atoms, which accounts for the optical thickness of the MOT. For the purposes of this experiment, it was sufficient to use the fluorescence imaging to provide a rough estimate of the number of atoms in the MOT. In subsequent stages of the experiment, where the atom number density was much lower, a photodiode with greater sensitivity than the CCD was used to measure the number of atoms. Details on this setup can be found in Section 6.4.1. Under the assumption that the power radiated per atom is constant, the power incident on the CCD P_{ccd} can be related to the scattering rate per atom as

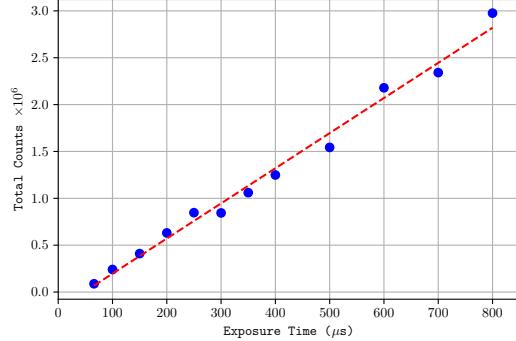


Figure 4.8: Integrated pixel counts as a function of CCD exposure time for an incident optical power of $0.17 \mu\text{W}$. The dashed line indicates a linear regression which gives a scaling factor of $\alpha = 2.2 \times 10^5 \text{ counts } \mu\text{s}^{-1} \mu\text{W}^{-1}$.

follows

$$P_{\text{ccd}} = \frac{\Omega}{4\pi} t R_{\text{sc}} \hbar \omega N_a \quad (4.1)$$

where $\Omega/4\pi = 1.8 \times 10^{-3}$ is the fractional solid angle subtended by the imaging optics, t is the transmission of the bandpass filter, R_{sc} is the scattering rate as previously defined in equation (??) and $\hbar \omega = 1.6 \text{ eV}$ is the emitted photon energy. This incident power is then related to the integrated number of pixel counts C_{int} by

$$C_{\text{int}} = \alpha \tau_{\text{exp}} \eta P_{\text{ccd}} \quad (4.2)$$

where τ_{exp} is the exposure time, $\eta = 0.14$ is the quantum efficiency of the CCD and α is a scaling factor that relates the total charge collected to the total number of pixel counts. By varying the exposure time used to image a collimated beam with a total power of $0.17 \mu\text{W}$, the total number of counts recorded by the camera as a function of exposure time is plotted in Figure 4.8. This gives a count scaling factor of $\alpha = 2.2 \times 10^5 \text{ counts } \mu\text{s}^{-1} \mu\text{W}^{-1}$.

4.3 Generating MOT light

All the MOT light in this experiment was generated by the μ Quans laser [4]. μ Quans is a French laser company that is a spin-off from the Institut d'Optique and Observatoire de Paris. A schematic of this laser system is shown in Figure 4.9. All of the light is fibre-coupled to minimise the number of free-space optical components and to make the system more stable in the presence of vibrations and small temperature variations. The μ Quans laser is comprised of four 1560 nm External-Cavity Diode Lasers (ECDLs) which are frequency-doubled to produce light at 780 nm. The first of these acts as a master laser which is locked to the $F = 3 \rightarrow F' = 3, 4$ crossover point in Rubidium-85 (^{85}Rb) and serves as an absolute frequency reference. The other three slave lasers are used for output. The first one is used to provide light for cooling, as well as repump light by modulating the phase of this laser using an Electro-optic Modulator (EOM). The other two make up a pair of lasers for driving Raman transitions. One laser is frequency-offset locked to the master and the other is phase-locked to the first, to ensure that the relative phase between the two lasers is constant. It should be noted that this Raman laser was not used in this experiment, so will not be discussed in further detail. Each of these slave lasers is amplified in an Erbium-Doped Fibre Amplifier (EDFA) before being frequency doubled in a Periodically Poled Lithium Niobate (PPLN) and passed through an AOM which is used to control the output power during the experiment.

4.3.1 Absolute Frequency Reference

The master laser is used to provide an absolute frequency to which the slave lasers are offset-locked. The frequency of the master is obtained using saturated absorption spectroscopy inside a Rubidium vapour cell. The sub-Doppler features in this spectrum

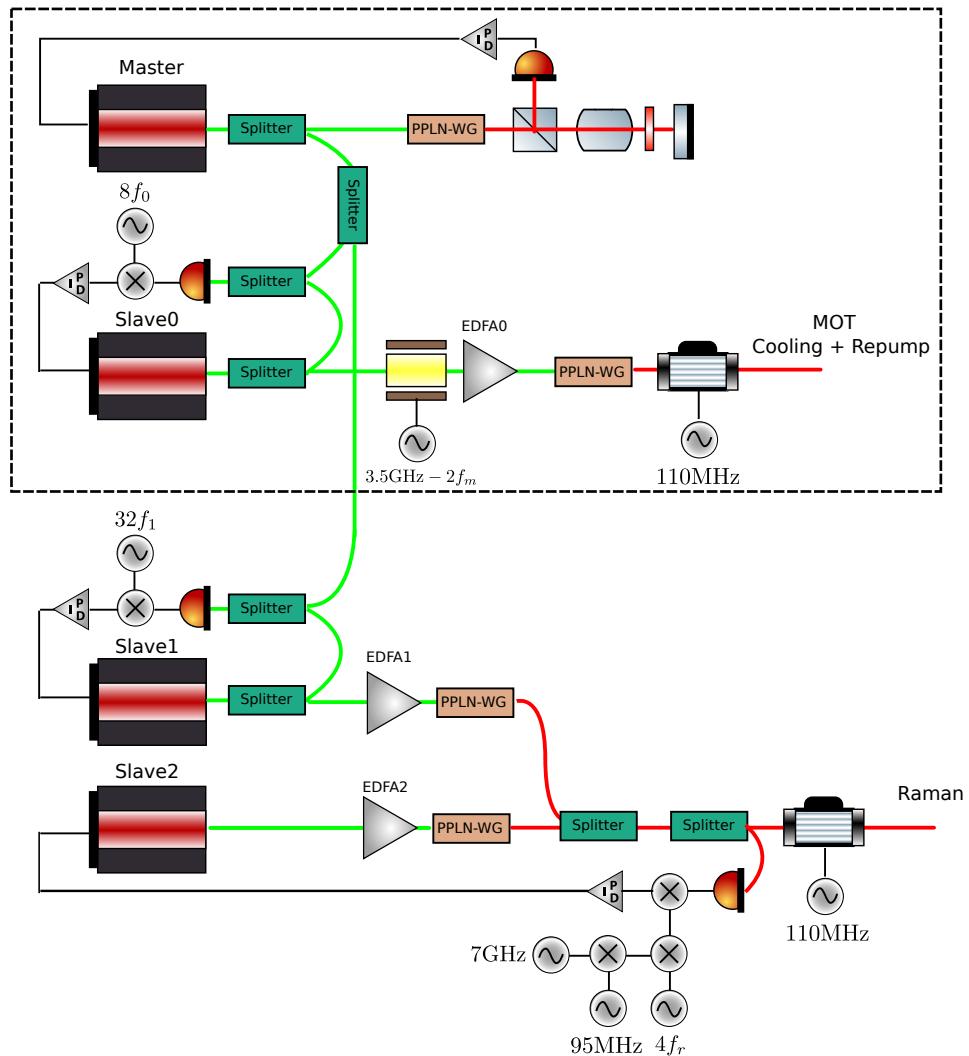


Figure 4.9: Schematic of the μ Quans laser system. Each output laser is derived from a 1560 nm ECDL (shown in green) which is amplified using an EDFA and then frequency-doubled to 780 nm using a PPLN crystal. A master laser is locked to the 3,4 crossover in ^{85}Rb and the output lasers are offset-locked to their corresponding frequencies. The dashed region indicates the components used for generating light for the MOTs, which was the only function of this laser for this experiment.

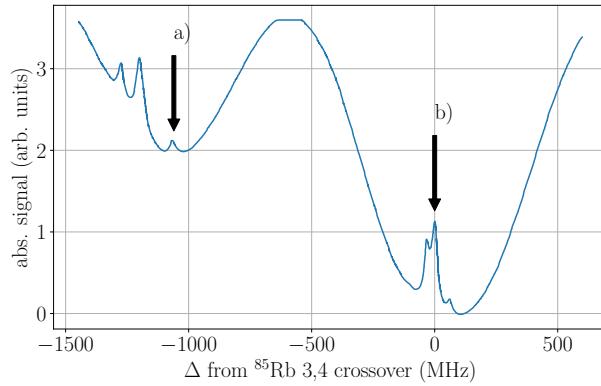


Figure 4.10: Saturated absorption spectroscopy using the Rubidium vapour cell in the μ Quans laser. The absorption features indicated are *a*: the $F = 2 \rightarrow F' = 3$ transition in ^{87}Rb and *b*: the crossover resonance between the $F = 3 \rightarrow F' = 3$ and $F = 3 \rightarrow F' = 4$ transitions in ^{85}Rb which is used to lock the frequency of the master laser.

are insensitive to temperature changes, and under sufficiently weak laser power have linewidths close to the natural linewidth of Rubidium ($\Gamma \sim 2\pi \times 6\text{MHz}$). Figure 4.10 shows the saturated absorption spectrum using the μ Quans master laser. This is obtained by fine adjustment of the temperature of the master **ECDL**. The master laser is set to lock to the crossover resonance between the $F = 3 \rightarrow F' = 3$ and $F = 3 \rightarrow F' = 4$ transitions in ^{85}Rb (indicated as *b*), which is the strongest feature in the spectrum. This absorption feature is around 1.1 GHz below the cooling transition in ^{87}Rb (indicated as *a*). The frequency of the laser is modulated to produce a signal which is proportional to the difference from the lock point. This is done by modulating the current to the master **ECDL**. The error signal shown in Figure 4.11 is obtained by demodulating the absorption signal using a lock-in amplifier. In addition to proportional feed-back from the error signal, the servo that controls the master laser frequency also contains an integrator to compensate for long-term drifts arising from temperature variations, for instance.

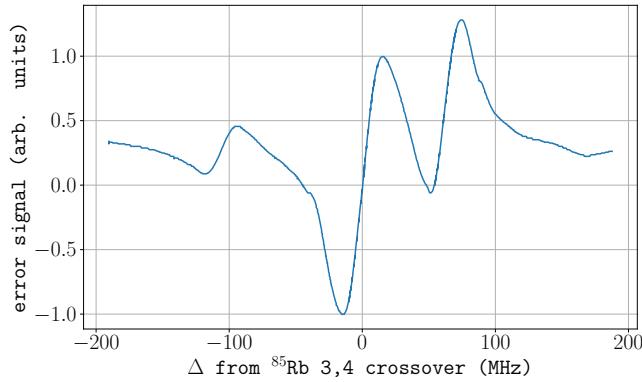


Figure 4.11: Error signal obtained by modulating the laser current. Close to the lock point, the signal is approximately linear. This signal is used in a feed-back loop to correct for frequency changes of the master laser.

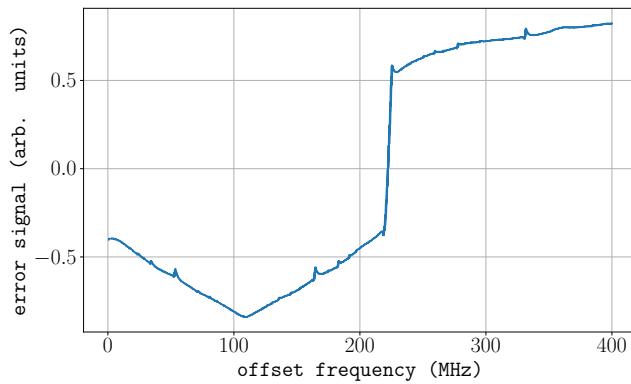


Figure 4.12: Error signal for the μ Quans cooling laser, plotted as a frequency difference from the start of the scan. This is obtained by comparing the beat frequency between the master and slave lasers to a reference frequency generated by a DDS. A servo loop feeds-back onto the frequency of the slave laser to keep this difference close to zero.

4.3.2 Cooling and Repump Light

The cooling light is generated by the first of the slave lasers. This is frequency-offset locked to the master by comparing their beat frequency to a local oscillator. The reference frequency is generated by a DDS and is scaled up by a factor of 8. A plot of the error signal used to lock this offset frequency is shown in Figure 4.12.

Light for driving the $F = 1 \rightarrow F' = 2$ repump transition is generated using an **EOM** that modulates the phase of the cooling laser to produce sidebands separated by integer multiples of the modulation frequency f_m . If the amplitude of the modulation frequency is small, only the first positive and negative sidebands are present. The modulation frequency is generated by another **DDS**, so that the frequency of the cooling and repump light can be ramped during the experiment independently (see Section 4.3.3). This is amplified, doubled and subtracted from a 3.5 GHz reference signal so that the positive frequency sideband is approximately 6.6 GHz above that of the cooling light, to address the repump transition. The power of this modulation signal can be externally controlled using a **vca!** (**vca!**) to control the ratio of repump power to cooling power. An RF switch is also used to switch the repump light on or off.

The total output power is controlled using an **AOM** that has a fixed modulation frequency of 110 MHz. Similarly to the repump **EOM**, a **vca!** is used to control the amplitude of this modulation to control the output power, as well as an RF switch to block the output.

4.3.3 Real-Time Control

During the experiment, it is necessary to vary the frequency and power of both the cooling and repump light. As mentioned above, the power is controlled using analogue and digital signals to the RF sources for the output **AOM** and **EOM**, respectively. Control of their frequency is achieved using the **DDS** that controls the corresponding reference frequency. In addition to updating the frequency of the slave or its sideband, the **DDS** can be programmed to ramp the frequency on a chosen output for a given duration and ramp rate. This is done during the experiment by sending serial messages to an application which interprets the command and programs each **DDS** using the

SPI interface. A glossary of the commands used is presented in Appendix ???. Once the **DDS** has been updated, the command is triggered to start using a digital signal so that the time at which the frequency updates is synchronised with the other hardware.

4.4 Controlling the MOTs

Effective trapping and cooling of Rubidium requires careful control of the light and magnetic fields used to create the **MOT**. Firstly, it is necessary to ensure the **MOT** beams are well polarised, with equal intensities at the centre of the trap. Otherwise, the imbalance of radiation pressure from each beam can exert a net force on the atoms, which causes the **MOT** to form at a position where the magnetic field is not zero [9]. When this field is switched off prior to cooling in an optical molasses, this force will push the atom cloud away from the equilibrium position and likely prevent further cooling. A similar argument holds in the case of balanced intensities and a stray magnetic field, which displaces the **MOT** from the centre of the beams to a position where the intensities are imbalanced. What follows is a description of the hardware used to provide more precise control of the relative power in each **MOT** beam, as well as the magnetic fields.

4.4.1 Optical Fibre Network

A network of fibre-based beam-splitters and **aoms!**s is used to distribute the light from the μ Quans fibre to each of the beams required for the 2D and 3D MOTs. This also provides a means of separate control of the power and frequency of the light at each output of the fibre network. A diagram of this setup is shown in Figure 4.13. The fibres are spliced together to minimise insertion loss and improve the alignment of the slow axes of connected fibres. Prior to this, the fibres were connected using mating sleeves,

which resulted in polarisation drifts that severely impacted on the stability of the MOTs. A set of *Gooch and Housego* fibre **aoms!**s with a central modulation frequency of 135 MHz are used to individually control the frequency and power of each output. The polarisation of the output from this fibre is purified using a polarising beam-splitter before a **HWP** aligns the polarisation axis of the light with the slow axis of a **PM** fibre. The light is first divided on a 1:2 beam-splitter, with 66% exiting one port, used for the 3D **MOT**. The 34% on the other port is then split again using another 1:2 beam-splitter so that 95% and 5% of the power exits each port to provide light for the 2D **MOT** and push beam, respectively.

The light for the 3D **MOT** is spilt using a 1:3 splitter into pairs of outputs for the light along the \vec{x} , \vec{y} and \vec{z} axes. Unlike the outputs along the other axes, the ones used for light along the \vec{z} axis have separate **aoms!**s. This is done so that during the experiment, a single beam along the \vec{z} axis can be used to blow away background atoms (see Section 5.3.4 for more details).

4.4.2 Magnetic Field Control

At different stages during the experiment, it is necessary to create a magnetic field to polarise the atoms along different axes. For this reason, the bias coils are all controlled using a bipolar current driver, where the output current is proportional to a control voltage. This control voltage is input at the non-inverting terminal of an OPA549 op-amp. The coils are placed in series with a sense resistor of resistance R_s at the output. The circuit is configured so that the voltage at the inverting terminal is $V_- = iR_s$. This forms a negative feed-back loop, so that if the output current drops due to ohmic heating of the coils, the op-amp increases the output voltage to keep the voltage at the two input terminals equal. The bias coils can be supplied with up to 1 A using a control voltage of 10 V.

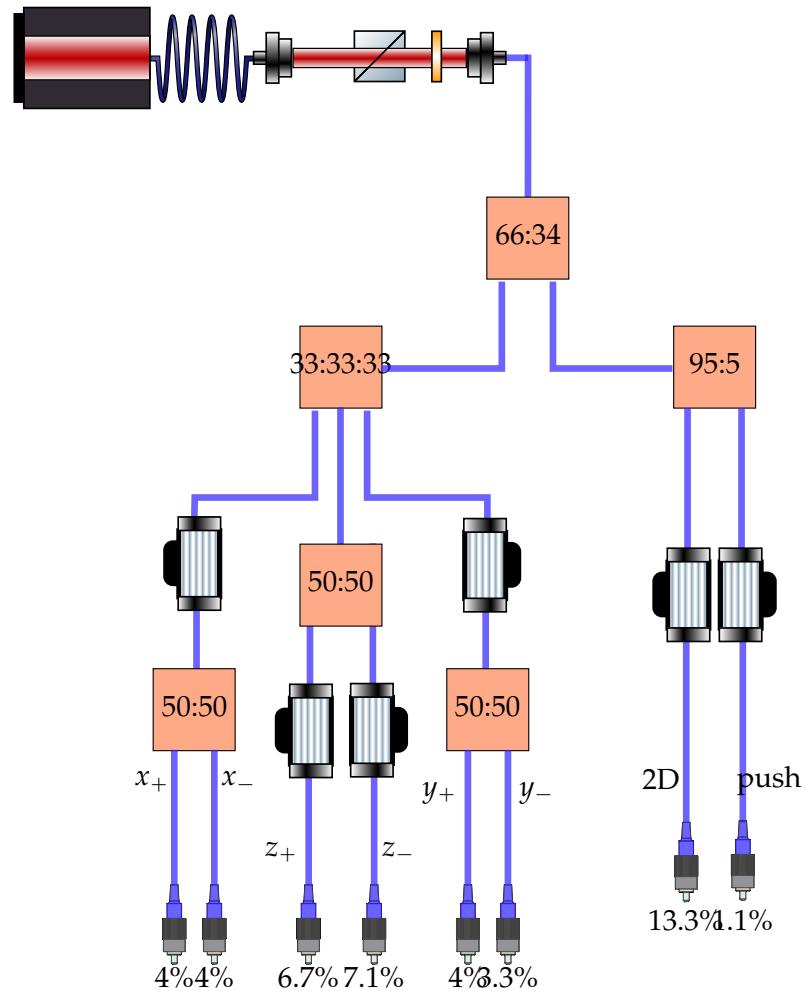


Figure 4.13: Fibre splitter and **aoms!**s for MOT! light distribution. The polarisation of the cooling and repump light from the μ Quans laser is aligned to the fibre network using a **PBS** and **HWP**. Apart from the outputs for the \vec{x} and \vec{y} 3D **MOT** beams, which have a single **AOM** per axis, the power and frequency at each output can be controlled independently. The percentages shown are the relative power at each output, accounting for insertion loss and driving each **AOM** with the optimum RF power.

This same circuit is used to control the 3D MOT coils, except for the fact that their larger resistance necessitates a larger gain to achieve the current necessary to produce a strong field gradient. During the experiment, the 3D MOT coils need to be switched off rapidly, to allow for effective sub-Doppler cooling of the atoms. Since the coils act as an inductive load, the time taken for the current is determined by the time constant $\tau = L/R$. However, if a negative voltage is applied, the energy stored in the inductor (and hence, magnetic field) will dissipate at a faster rate. To determine the minimum time taken for the field to switch off, a flux-gate magnetometer was used to measure the time taken for the magnetic field outside the chamber to reach a steady value after changing the voltage across the coils. With an applied voltage of 0 V, a $1/e$ decay time of 2.5 ms. Under the maximum voltage of -24 V that the current driver can output, the field can be completely switched off in 800 μ s.

The quadrupole coils for the 2D MOT are controlled in a different manner. They are switched off using an IGBT to cut the flow of current when the gate voltage drops below a threshold value. This generates a large back-EMF that opposes the changing current. To prevent damage to the transistor, a diode and 10Ω power resistor are placed in parallel with the coils. This allows the current generated by the back-EMF to dissipate without damaging the IGBT. These field from these coils can be switched off in less than 1 ms, but this is less critical than the time taken for the 3D MOT quadrupole field to switch off.

4.5 Characterising the MOTs

This section discusses the performance of the 2D and 3D MOTs for trapping and cooling ^{87}Rb . The main goal of this stage of the experiment is to quickly produce an ensemble of trapped, cold atoms in the 3D MOT. For this reason, the loading rate of

2D MOT				3D MOT			
	Laser Power	Magnetic Field		Laser Power	Magnetic Field		
Cooling	60 mW	$d\vec{B}/d\rho$	18 G cm^{-1}	Cooling	130 mW	$d\vec{B}/dz$	15 G cm^{-1}
Repump	6 mW	B_x	0.48 G	Repump	13 mW	B_x	1 G
Push	500 μW	B_y	-0.46 G			B_y	-0.5 G
						B_z	0.22 G

Table 4.2: Typical optical and magnetic parameters used for the 2D and 3D MOTs. The optical powers listed are the total used for each MOT, which is divided into separate beams. The bias field strengths are the values used during the preliminary trapping stage of the experiment. The specified field gradients are given along the radial direction and the symmetry axis of the quadrupole coils for the 2D and 3D MOTs, respectively.

the 3D MOT is a useful figure-of-merit. As further cooling in an optical molasses is necessary to achieve a sufficiently cold ensemble for interferometry (see Section 5.2), the temperature of atoms in the MOT will not be discussed in detail.

At the start of the experiment, the light and magnetic fields to produce the 2D and 3D MOTs are switched on. Table 4.2 shows the typical values for the cooling and repump power, as well as the field gradients and bias fields used. The cooling light is detuned by -2Γ from the $|F = 2\rangle \rightarrow |F' = 3\rangle$ transition for the 2D MOT and -2.5Γ for the 3D MOT, whereas the push beam is at resonance. The light and magnetic fields for the 2D MOT are switched off after 100 ms and the 3D MOT is kept on for a further 50 ms to allow for the transit of the remaining atoms from the 2D MOT to the 3D MOT. After a sufficient number of atoms are loaded, the experiment proceeds by switching off the 3D quadrupole field prior to cooling in an optical molasses.

4.5.1 3D MOT Loading Rate

The loading rate of the 3D MOT from a beam of atoms originating from the 2D MOT can be understood using the following rate equation

$$\frac{dN}{dt} = R\phi_{rb} - (\alpha\phi_{rb} + \beta n_{bg}) N - \gamma N^2 \quad (4.3)$$

where ϕ_{rb} is the flux of rubidium through 3D MOT capture volume and R describes the rate at which rubidium is cooled and trapped such that $R\phi_{rb}$ is the loading rate of the 3D MOT. The second term describes a loss rate due to collisions between trapped atoms and untrapped rubidium and background atoms. These loss rates are parameterised by α and β , respectively. The final term describes the loss of atoms from the trap due to intra-trap collisions [7] which depends on the density of atoms in the trap. In the case of a large flux of atoms from the 2D MOT the first two terms dominate, leading to a simple solution for the number of atoms in the 3D MOT

$$N(t) = \frac{R\phi_{rb} \left(1 - e^{-t(\beta n_{bg} + \alpha\phi_{rb})}\right)}{\beta n_{bg} + \alpha\phi_{rb}} \quad (4.4)$$

which has a steady-state atom number given by

$$N_\infty = \frac{R\phi_{rb}}{\beta n_{bg} + \alpha\phi_{rb}} \quad (4.5)$$

Under a small atomic flux, both the loading rate and steady-state atom number increase as the flux of atoms from the 2D MOT increases. Once this flux is great enough, the loss due to background atom collisions is small compared to the loss due to rubidium collisions and the final number is independent of ϕ_{rb} .

To optimise the 3D MOT loading rate, we varied both the . The flux of atoms from the 2D MOT depends on the 2D MOT loading rate, which in turn depends on the capture

volume and rubidium number density inside the cell. Without any longitudinal cooling, a significant fraction of atoms leaving the cell will have a velocity greater than the 3D MOT capture velocity [8].

The effect of varying the partial pressure of rubidium inside the source cell are shown in Figure 4.14. For a range of dispenser currents, the number of atoms as a function of loading time was measured. At low partial pressures, the loading rate increases due to the increase in the flux from the 2D MOT. As the pressure increases, the increasing flux gives a larger steady-state number of atoms, up until the background pressure becomes negligible. Figure ?? and Figure ?? compare the loading curves observed with and without the push beam. Below a threshold pressure, the push beam greatly improves the loading rate since a greater fraction of the atoms can be captured in the 3D MOT. Figure 4.14c shows the fitted loading rate for each scenario and indicates a clear optimum pressure, where the loading rate is maximised and the same steady-state atom number is reached with or without the push beam. Above this pressure, the loading rate is sharply reduced as the increased collision rate between cold atoms from the 2D MOT and hot untrapped ones reduces the atomic flux. This also increases the mean velocity of atoms in the beam, since faster ones are less likely to collide with a background atom before exiting the cell. At very high pressures, the mean velocity is so great that only a small fraction of atoms can be captured and the push beam has little effect on the loading rate.

4.6 Conclusion

This chapter has introduced the components of the experiment that were used to trap and cool atoms in a MOT. This is used to prepare an ensemble of cold atoms in a pure quantum state, suitable for interferometry. An optimisation of the loading rate of the

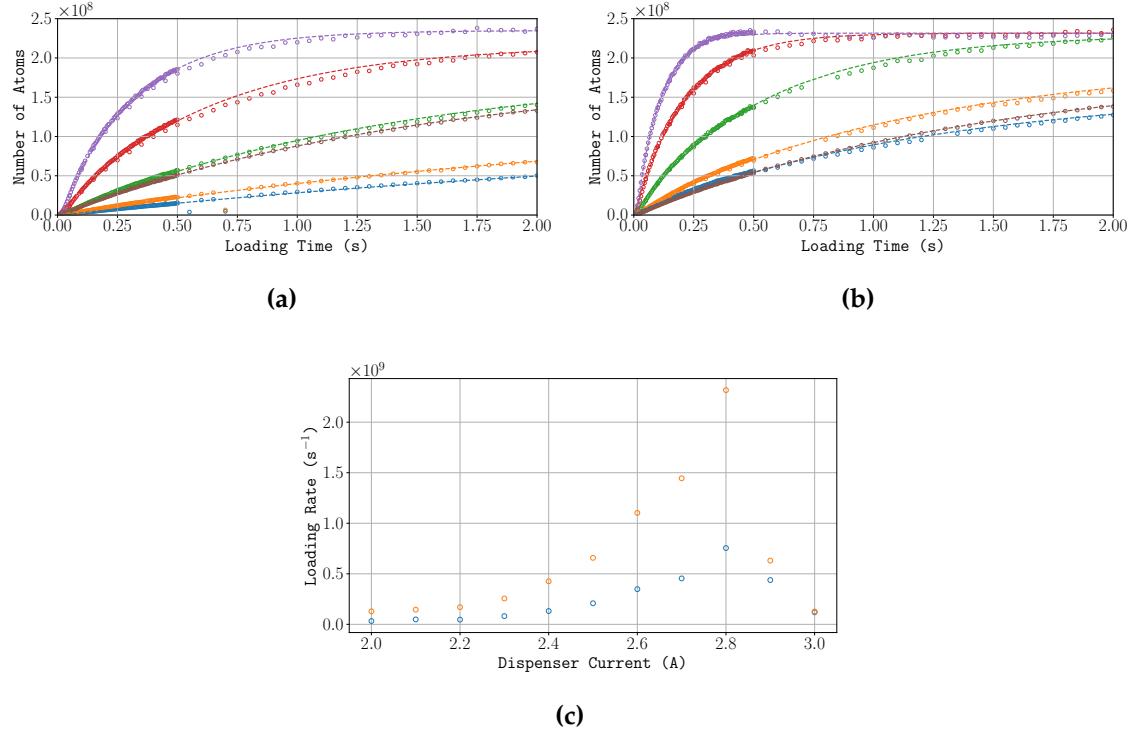


Figure 4.14: Number of atoms in the 3D MOT over time for a range of dispenser currents. For clarity, only the loading curves for dispenser currents of 2 A (blue), 2.2 A (orange), 2.4 A (green), 2.6 A (red), 2.8 A (purple) and 3 A (brown) are shown in (a) and (b), which present the number of atoms over time without and with a push beam, respectively. The loading rate ($R\phi_{rb}$ in equation (4.3)) in both instances is shown in (c). As the partial pressure of rubidium increases, the flux of atoms from the source cell increases. By longitudinally cooling the atoms, the push beam enhances the loading rate of the 3D MOT. Above a dispenser current of 2.8 A, the collision rate with hot untrapped atoms greatly reduces the atom flux, reducing both the loading rate and steady-state atom number.

3D MOT was carried out to reduce the dead time between consecutive measurements of acceleration using the interferometer.

Chapter 5

Preparing Atoms for Interferometry

5.1 Chapter Overview

This chapter presents the stages of the experiment which prepare an ensemble of atoms for interferometry, after they are loaded into the 3D MOT. After being released from the trap, the atoms are cooled and launched using a moving molasses, as described in Section 5.2. Following this, a sequence of optical and microwaves pulses are used to increase the population in the $|1, 0\rangle$ ground state and end with an ensemble which has a narrow velocity spread along the Raman axis. A characterisation of this is given in Section 5.3.

Some sections of this chapter refer to parts of the experiment which have yet to be introduced. Details on the Raman laser and the velocity-selective Raman pulse can be found in Section 6.3.1 and Section ??, respectively. A description of the detection scheme, used to measure the population of atoms in $|F = 1\rangle$ and $|F = 2\rangle$ is presented in Section 6.4

5.2 Cooling in Optical Molasses

Introduce motivation here A low thermal velocity means that the atoms can be interrogated for a longer time and in the case of atom interferometry, leads to a more sensitive measurement of acceleration. In addition, the thermal expansion of the ensemble leads to greater systematic phase shifts due to effects such as magnetic field gradients and laser wavefront distortions. The temperature of atoms inside the MOT is greater than desired, so further cooling is required before a strong interferometer signal can be achieved. Temperatures well below that of the Doppler limit ($146 \mu\text{K}$ for ^{87}Rb) can be reached using the dissipative force that acts on an atom travelling through a spatially varying electric field.

This section describes the work towards to cooling and launching the atoms in a moving optical molasses. It starts with a motivation for launching the atoms in Section 5.2.1. The following section discusses the control of the intensity and frequency of the light during the molasses stage of the experiment. A description of the techniques needed to cool the atoms in a moving molasses is then given in Section 5.2.3. Finally, this section concludes in Section 5.2.4 with measurements of both the temperature and trajectory of the atom cloud which were measured using a ballistic expansion method.

5.2.1 Motivation for Launching

As previously discussed in Section 2.5.1, there are two pairs of counter-propagating beams which can drive Raman transitions between the two hyperfine ground states. If an atom can be stimulated by both pairs, then the additional trajectories this introduces do not interfere, resulting in a reduction in the fringe visibility. This problem can be avoided by using the fact that the Raman transition is Doppler-sensitive to ensure that the atoms are only driven by one pair of beams. Each pair has an opposite Doppler shift

$\pm\omega_D = \pm\mathbf{k}_{\text{eff}} \cdot \mathbf{v}$ and so their transition frequencies are separated by $2\omega_D$. Therefore, the atoms are launched so that their centre-of-mass velocity along the Raman axis is large enough to lift the degeneracy of the two Raman transitions.

5.2.2 Frequency and Power Control

A timing diagram illustrating the power and frequency during this stage of the experiment is shown in Figure 5.1. After the atoms are loaded into the MOT, they are released by switching off the quadrupole field. Once this field has decayed away, the frequency and intensity of the cooling light are ramped adiabatically. The frequency of the cooling light is ramped to -25Γ over 1.4 ms. Since the repump light is generated using an EOM, the modulation frequency is simultaneously ramped up to keep this light resonant with the $|F = 1\rangle \rightarrow |F' = 2\rangle$ transition. Additionally, the relative detuning of counter-propagating MOT beams is varied so that the atoms are cooled into a moving molasses (see Section 5.2.3). After this, the intensity of the light is reduced over 5 ms. The response of the output AOM on the μ Quans laser was calibrated so that we could apply a voltage ramp that gives an approximately linear intensity ramp.

5.2.3 Launching in a Moving Molasses

The configuration for launching atoms along the Raman axis is shown in Figure 5.2. The forward-propagating beams are blue-detuned by $+\delta_l$ and the backward-propagating ones are red-detuned by $-\delta_l$, so that atoms with a velocity along the beam axis of $\vec{v} = \delta_l \lambda$ are+ resonant with both beams. The frequency of each beam is ramped from the initial value by varying the modulation frequency of its AOM. This ramp occurs slowly to ensure the atoms are adiabatically accelerated to the resonant velocity, minimising excess heating of the atoms. As there is no pair of MOT beams along the axis

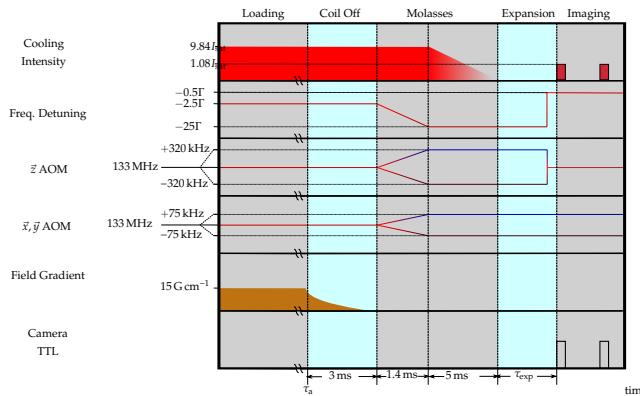


Figure 5.1: Timing diagram for the molasses stage of the experiment. After a time $\tau_a = 100$ ms, the atoms are released from the MOT. The molasses sequence begins 3 ms later, once the magnetic field from the MOT coils has decayed away. First, the frequency of the cooling light is ramped to -25Γ over 1.4 ms. The relative frequencies of counter-propagating MOT beams are detuned so that the atoms are cooled in a moving frame, launching them along a parabolic path (see Section 5.2.3). Next, the intensity of the MOT light is reduced linearly over 5 ms. To measure the temperature, the atoms are left to expand for a duration of τ_{exp} ms, after which they are imaged using the camera.

of the Raman beams, the \vec{x} and \vec{y} MOT beams, whose axes are nominally at 45° to the Raman axis, are used to launch the atoms. By controlling the power and alignment of each beam, the net velocity on the atoms will be along the Raman axis. If the detuning of both pairs of beams is the same, then the velocity along the Raman axis is given by $\vec{v}_r = \sqrt{2}\delta_l \lambda$.

As well as launching the atoms horizontally, the atoms are launched vertically so that they do not fall as far under gravity. Since the atoms remain close to the centre of the beam, where the intensity across the cloud is more uniform, longer pulse separation times can be used before the intensity gradient causes a significant dephasing. This launch is carried out using the MOT beams that lie along the vertical \vec{z} axis. Under an appropriate choice of launch velocity, the centre-of-mass trajectory transverse to the Raman beam wavefront can be chosen to maximise the interferometer fringe visibility. This is presented in further detail in Section 6.6.1.

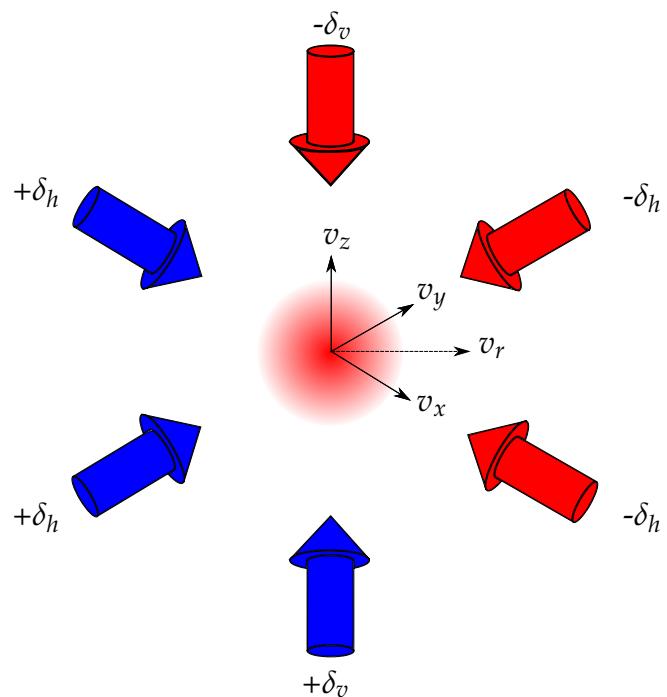


Figure 5.2: Beam configuration for a moving molasses. When counter-propagating beams are detuned from each other, the atoms are slowed to a velocity which balances the frequency of each beam. Along the vertical axis, the beams are detuned by δ_v so that the cloud is launched upwards with a velocity $v_z = \delta_v \lambda$. In the horizontal plane, the \vec{x} and \vec{y} beams are detuned by δ_h so that the resultant velocity is along the Raman axis $v_r = \sqrt{2} \delta_h \lambda$.

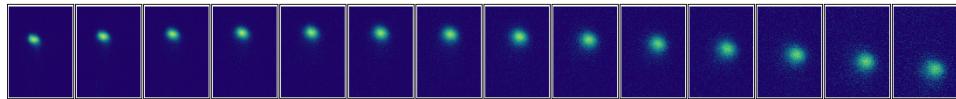


Figure 5.3: A series of images showing the trajectory of the atom cloud after being cooled in a moving molasses. The first image was taken 7 ms after initiating the molasses and a subsequent one every 5 ms. Each image represents a region of interest of dimensions 1150×1650 pixels that covers the spatial extent of the atom cloud during the launch.

5.2.4 Imaging the Atom Cloud over Time

Once released from the trap, the atom cloud is free to expand. Provided that there are no external forces on the atoms from electric or magnetic fields, the expansion of the cloud is determined by its thermal velocity distribution. In addition to this, the centre-of-mass moves due to its initial velocity and acceleration due to gravity. For the purposes of this discussion, it is worthwhile to consider the motion of atoms within the ensemble separately to the centre-of-mass motion as these provide a means of measuring the temperature and launch velocity, respectively. These were measured by imaging the distribution of atoms after allowing the cloud to expand in the dark for a range of expansion times. A typical atom cloud trajectory is shown in Figure ??, in which the cloud was imaged up to 76 ms after being released from the trap.

Measuring the Temperature

In thermal equilibrium, the velocity distribution of the atoms is described by a Maxwell-Boltzmann distribution

$$f(v_x, v_y, v_z) = \left(\frac{m}{2\pi k_B} \right)^{3/2} e^{-\frac{m(v_x^2 + v_y^2 + v_z^2)}{2k_B T}} \quad (5.1)$$

where $v'_i = v_i - \langle v_i \rangle$ is the difference from the average velocity. Along a single axis, the velocity distribution is obtained by integrating over the other velocity components.

For the sake of notation, the following discussion uses x and v_x as labels for position and velocity, but these are interchangeable with the equivalent components along the other axes. As equation (5.1) is a product of velocity distributions along each axis, the velocity distribution along one axis is

$$f(v_x) = \left(\frac{m}{2\pi k_B} \right)^{1/2} e^{-\frac{m(v_x - \langle v_x \rangle)^2}{2k_B T}} \quad (5.2)$$

Suppose that there are initially $n_0(x)dx$ atoms within the region $(x, x + dx)$, where $n_0(x) = n(x, t = t_0)$ is the initial atomic density along one axis. During ballistic expansion, the atoms redistribute themselves according to their velocity. After a time t of free expansion, the position of an atom initially at x is $x + v_x t$. Assuming that the number density is initially a Gaussian, with a peak number density $n_0(x_0)$ at the centre-of-mass, the number density at later times is given by a convolution with equation (??)

$$n(x, t) = \int n_0(x_0) \left(\frac{m}{2\pi k_B} \right)^{1/2} e^{-\frac{m(v_x - \langle v_x \rangle)^2}{2k_B T}} e^{-\frac{(x + v_x t - x_0)^2}{2\sigma_0^2}} dv_x \quad (5.3)$$

where σ_0 is the $1/e^2$ initial width of the cloud. As a convolution of two Gaussians, equation (5.3) is also a Gaussian, with a $1/e^2$ width given by

$$\sigma(t)^2 = \sigma_0^2 + \frac{k_B T}{m} t^2 \quad (5.4)$$

Figure 5.4 shows the measured cloud width over a range of expansion times. The inset shows a typical density profile along each axis, obtained by imaging the cloud on a camera, as previously described in Section ???. For the purposes of measuring the temperature, the total atom number and the initial cloud size are not important, so no attempt was made to estimate these. The initial measurement was made 7 ms after the end of the molasses to allow for enough time to re-lock the laser to -0.5Γ below

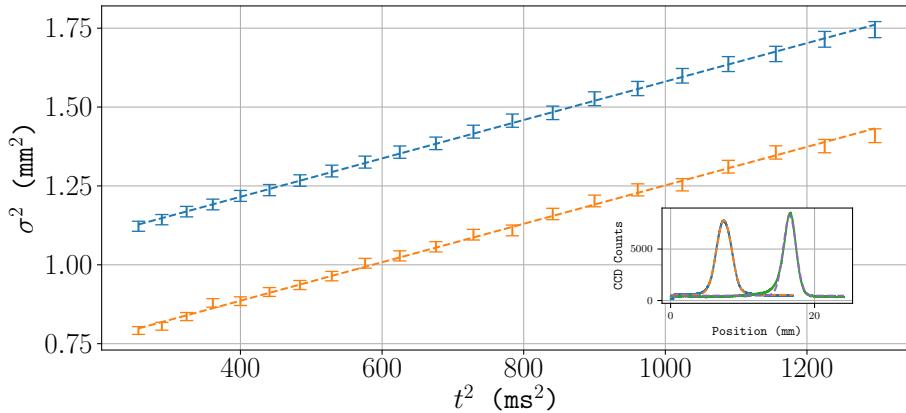


Figure 5.4: Atom cloud temperature using a ballistic expansion measurement. After molasses, the cloud is left to expand in the dark and in a region of close to zero magnetic field. A least-squares linear fit of $\sigma(t)^2$ is used to estimate the temperature using equation (??). The inset shows a typical density profile from a single image, obtained by integrating the signal from a CCD camera along the two axes of the sensor. The gradient from the fits for the horizontal (blue) and vertical (orange) axes are $T_x = 6.38(15)$ µK and $T_y = 6.38(12)$ µK.

the $|F = 2\rangle \rightarrow |F' = 3\rangle$ transition and align the bias field to the \vec{z} axis so that the atoms could be optically pumped into the $|2, 2\rangle$ state. At each measurement time, a non-linear least squares fit to equation (5.3) along each axis was carried out to estimate the width of the cloud. Then, a least squares linear fit was used to estimate the temperature along each axis from the gradient, as per equation (5.4). The measured temperature along the horizontal and vertical axes of the camera was $T_x = 6.37(15)$ µK and $T_y = 6.38(12)$ µK, respectively.

Measuring the Launch Trajectory

The same method used to measure the temperature of the cloud can also be used to measure the position of the centre-of-mass. In this case, the quantity of interest is $\langle x(t) \rangle$. Since the cloud is in free-fall, the trajectory for the centre-of-mass is then given by the well-known equation-of-motion for a particle moving under constant

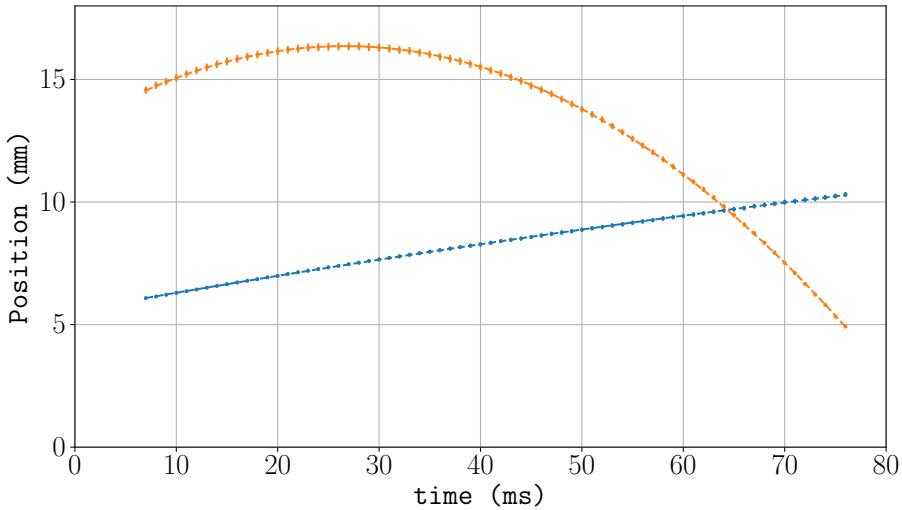


Figure 5.5: Measured centre-of-mass position over time. The horizontal component of the position is shown in blue and the vertical in orange. Each trajectory is fit to equation (5.5) to estimate the launch velocity. The best-fit values are $v_v = 25.00(35) \text{ cm s}^{-1}$ and $a_v = -9.400(75) \text{ m s}^{-2}$ along the vertical axis and $v_h = 7.39(21) \text{ cm s}^{-1}$ and $a_h = -0.310(52) \text{ m s}^{-2}$ along the horizontal.

acceleration

$$\langle x(t) \rangle = \langle x(0) \rangle + v_x t + \frac{1}{2} a_x t^2 \quad (5.5)$$

where v_i is the initial velocity along the given axis and a_i is the acceleration.

To launch the atoms both vertically and horizontally (along the axis parallel with the Raman light), the (z_+, z_-) **aoms!**s were ramped so that the relative frequency difference between each beam was $2 \times 320 \text{ kHz}$ and the x and y **AOM** frequencies were ramped to give a frequency difference of $2 \times 75 \text{ kHz}$ between the horizontal **MOT** beams. Figure 5.3 is a plot of the measured centre-of-mass position along the horizontal and vertical camera axes over time. A linear least-squares fit to equation (5.5) gives a vertical launch quantities of $v_v = 25.00(35) \text{ cm s}^{-1}$ and $a_v = -9.400(75) \text{ m s}^{-2}$ and $v_h = 7.39(21) \text{ cm s}^{-1}$ and $a_h = -0.310(52) \text{ m s}^{-2}$ along the horizontal axis.

When compared to the expected velocities from the detunings, $v_v^{(l)} = 24.96 \text{ cm s}^{-1}$ and $v_h^{(l)} = 5.85 \text{ cm s}^{-1}$, the measured horizontal velocity is far greater than expected. This can be explained by a residual magnetic field, that is not cancelled using the bias coils. In the presence of a magnetic field, atoms cooled in an optical molasses are decelerated to a velocity at which the Zeeman shift is cancelled by the Doppler shift. This Velocity-selective resonance (**VSR**) depends on the orientation of the magnetic field to the polarisation of the light. In a one-dimensional optical molasses, a vsr occurs at $v_{\text{res}}^{(1)} = -\mu_B g_F B / \hbar k$ when the magnetic field is aligned with the wave-vector of the light [VanderStraten1993]. When the field is aligned at an arbitrary angle an additional resonance at $v_{\text{res}}^{(2)} = -\mu_B g_F B / 2\hbar k$ is present, due to additional $(\sigma^\pm - \pi)$ transitions [Chang2002]. A residual field along the Raman axis of 20 mG would shift the **VSR** along \vec{x} and \vec{y} by 1.09 cm s^{-1} corresponding to a velocity of 1.54 cm s^{-1} along the Raman axis.

5.3 State Preparation

After the atoms have been cooled in an optical molasses, the population will mostly be distributed across the $|F = 2\rangle$ level, along with a small fraction distributed across the $|F = 1\rangle$ level. The Raman transition only couples the $|1, 0\rangle$ and $|2, 0\rangle$ states, so atoms in the other hyperfine ground states cannot participate in the interferometer. In fact, since the individual Zeeman sub-levels are not resolved during detection, these background atoms result in a loss of fringe visibility. One way to overcome this is to apply a pulse of light resonant with the $|F = 2\rangle \rightarrow |F' = 3\rangle$ transition to push the non-participating atoms out of the interferometer detection region. Of course, this must be done after applying a Raman pulse to transfer an ensemble of atoms from $|2, 0\rangle$ into $|1, 0\rangle$. In this simple scheme a large fraction of the atoms are removed, which is undesirable since

measurements of a low number of atoms are inherently more uncertain due to shot number fluctuations.

The following section discusses a method of preparing the atoms to increase the population in the $|1, 0\rangle$ ground state. An overview of the scheme is given in Section 5.3.1. This is followed by a discussion of the initial steps which optically pump atoms into the $|1, 0\rangle$ state in Section 5.3.2. A description of the microwave pulse used to drive atoms into the $|F = 2\rangle$ level is given in Section ???. This section concludes with the method used to blow away the atoms which do not contribute to the interferometer in Section 5.3.4. A key step which has been omitted is the velocity-selective Raman pulse. This is described in more detail later, in Section ???

5.3.1 Schemes for Preparation

The scheme used to prepare atoms in the $|1, 0\rangle$ state is the following:

1. Light resonant with the $|F = 2\rangle \rightarrow |F' = 2\rangle$ transition pumps atoms into the $|F = 1\rangle$ level
2. Light resonant with the $|F = 1\rangle \rightarrow |F' = 0\rangle$ transition drives (σ^\pm) transitions to pump atoms into the $|1, 0\rangle$ dark state
3. A microwave π -pulse transfers atoms to $|2, 0\rangle$
4. A Raman π -pulse transfers atoms with a narrow velocity spread back to $|1, 0\rangle$
5. The atoms which remain in $|F = 2\rangle$ are blown away

A diagram of the population of each hyperfine ground state and the laser frequencies used to drive these transitions is given in Figure 5.6. With the exception of step 4, the light is provided by the μ Quans laser using the MOT collimators aligned to the vertical \vec{z} axis. The frequency of the cooling laser and the repump sideband are set so

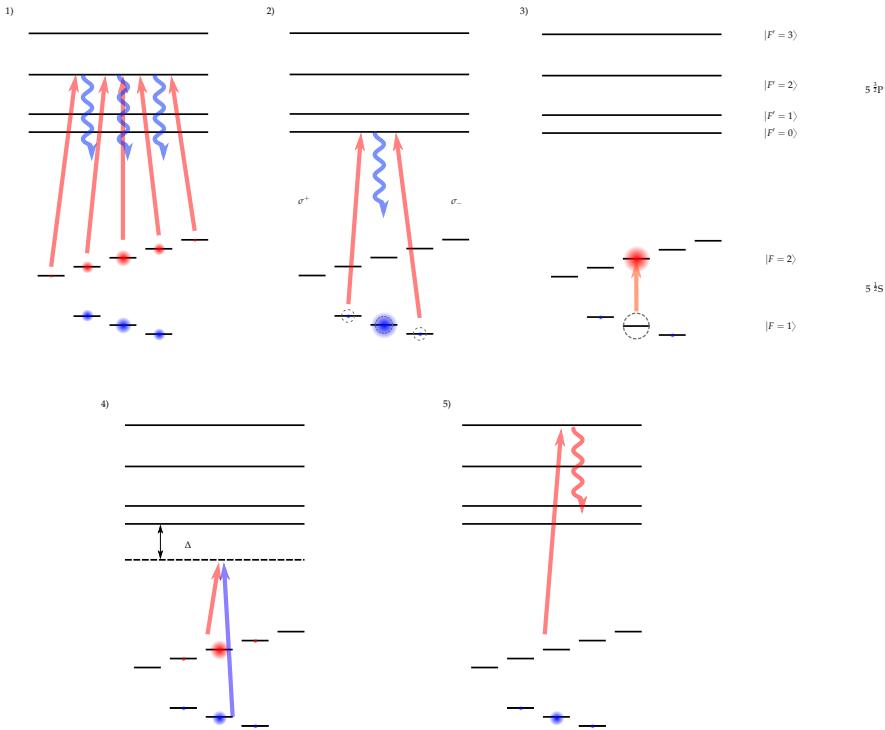


Figure 5.6: Sequence of optical and microwave pulses used to prepare an ensemble of atoms in $|1,0\rangle$. The red arrows indicate optical transitions to and from $|F = 2\rangle$ and equivalently for the blue arrows and $|F = 1\rangle$. A residual population in the $|1, \pm 1\rangle$ states is present, which contributes to a background during the interferometer.

that the relevant transitions for steps 1 and 2 are addressed. As the $|F = 1\rangle$ light is a sideband of the $|F = 2\rangle$ light, it is not possible to blow away atoms in $|F = 1\rangle$ without also blowing away atoms in $|F = 2\rangle$. This problem is overcome by using microwave pulses to drive atoms up to $|F = 2\rangle$ before velocity selection.

A timing diagram of the state preparation sequence is shown in Figure 5.7, which indicates the duration for which each optical or microwave pulse is applied, as well as the direction of the applied magnetic field. The field is switched slowly over 2 ms (which is omitted from the diagram) to preserve the spin state of each atom.

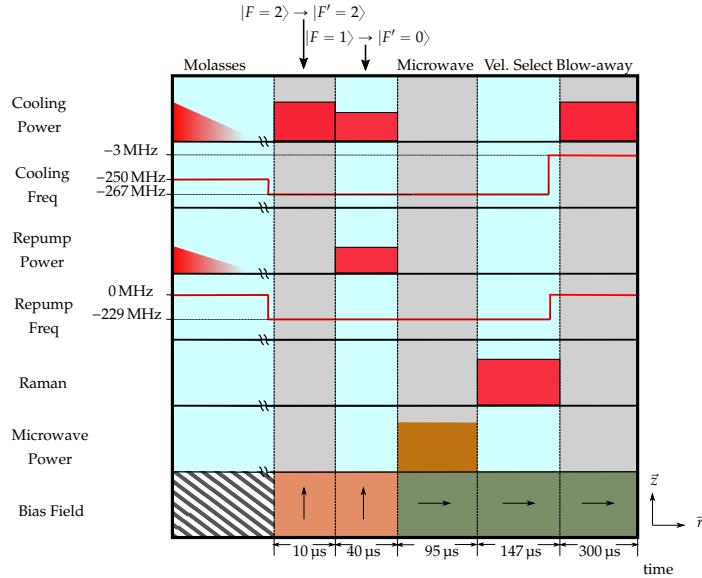


Figure 5.7: Timing diagram for state selection sequence. The durations labelled are indicative of the time required to drive the atoms into the desired state at each step. After the $|F = 1\rangle \rightarrow |F' = 0\rangle$ pumping, the magnetic field is re-oriented along the Raman axis \vec{r} . The 2 ms field switching time has been omitted.

5.3.2 Optically Pumping the Atoms

Driving the $|F = 2\rangle \rightarrow |F' = 2\rangle$ transition

After the molasses, the frequency of cooling light is 150 MHz below the $|F = 2\rangle \rightarrow |F' = 3\rangle$ transition. This can off-resonantly excite an atom to the $|F' = 2\rangle$ excited level, but the small scattering rate means that on average, an atom will need to scatter many photons before it is pumped into the $|F = 1\rangle$ level. Therefore, to minimise the heating during this pumping process, the frequency of the cooling light is resonant with the $|F = 2\rangle \rightarrow |F' = 2\rangle$ transition.

Figure 5.8 shows the population in the two hyperfine ground states as the duration of the $|F = 2\rangle \rightarrow |F' = 2\rangle$ light is increased. The rate at which atoms are pumped into $|F = 1\rangle$ increases with the strength of the applied magnetic field. At zero field, there exists a dark state which is a coherent superposition of the $|2, m_F\rangle$ states [1]. Applying a magnetic field lifts the degeneracy between the Zeeman sub-levels so that this dark

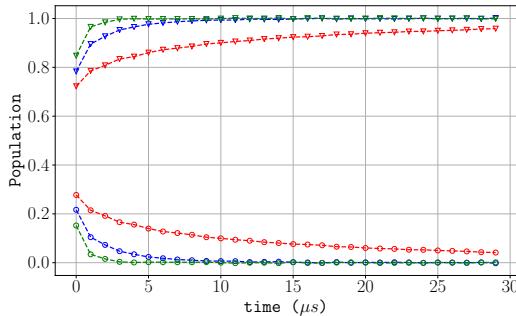


Figure 5.8: Population across the two hyperfine ground states after $|F = 2\rangle \rightarrow |F' = 2\rangle$ pumping under various magnetic field strengths. The \triangledown (\circ) markers indicate the population in $|F = 2\rangle$ ($|F = 1\rangle$). The red, blue and green series correspond to field strengths of -0.16 G , 1.67 G , and 3 G , respectively.

state is no longer stationary. The evolution rate of this dark state, and hence pumping rate, increases with an increasing Zeeman shift. At a field strength of 3 G , the atoms can be pumped into $|F = 1\rangle$ in less than $5\text{ }\mu\text{s}$.

Driving the $|F = 1\rangle \rightarrow |F' = 0\rangle$ transition

After this first pumping step, the atoms are distributed across the Zeeman sub-levels in $|F = 1\rangle$. The next pulse of light is used to increase the population in $|1, 0\rangle$ by driving $|F = 1\rangle \rightarrow |F' = 0\rangle$ transitions. During this time, the $|F = 2\rangle \rightarrow |F' = 2\rangle$ light remains on which helps to prevent atoms from populating the $|F = 2\rangle$ level through off-resonant $|F = 1\rangle \rightarrow |F' = 1\rangle$ excitations. The magnetic field present means that the circularly-polarised \vec{z} MOT beams only drive σ^\pm transitions, so the $|1, 0\rangle$ state is in principle a dark state.

The distribution of atoms across the Zeeman sublevels was measured using a microwave pulse to drive atoms into the $|F = 2\rangle$ level, which is described in Section 5.3.3. For each π microwave transition, the frequency of the microwave field was varied to find the resonant frequency. The resulting spectra for $m_F = -1$ and $m_F = 0$ are shown in Figure 5.9, both with and without applying light to pump into the $|1, 0\rangle$ state.

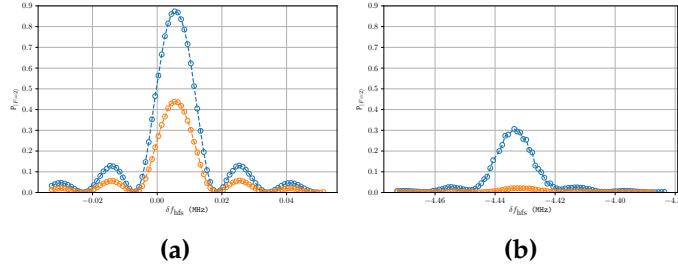


Figure 5.9: Population of atoms in (a) $|1, 0\rangle$ and (b) $|1, -1\rangle$, measured by applying a $68\ \mu\text{s}$ microwave pulse to drive atoms into the $|F = 2\rangle$ level. The orange and blue points indicate the measured populations with and without $|F = 1\rangle \rightarrow |F' = 0\rangle$ pumping. The microwave frequency is plotted as a detuning from the hyperfine splitting frequency f_{hfs} .

The $0 \rightarrow 0$ clock transition is detuned from the hyperfine splitting frequency due to the applied magnetic field, giving a second-order Zeeman shift of $515\ \text{Hz}\ \text{G}^{-2}$. The measured shift of $5.6\ \text{kHz}$ corresponds to a field strength of $3.3\ \text{G}$.

A plot of the population in each Zeeman sub-level for increasing pumping times is given in Figure 5.10. In this instance, the optical pumping does not completely deplete the population from the $m_F = \pm 1$ sub-levels. After pumping for $30\ \mu\text{s}$, approximately 5% of the population remains in the $m_F = \pm 1$ sub-levels. The $|1, 0\rangle$ state can only be excited to $|F' = 0\rangle$ by π -polarised light, which suggests that the magnetic field is mis-aligned with the \vec{z} MOT beams. The effect of these background atoms on the measured interferometer signal is discussed later, in Section 6.4.2.

5.3.3 Including Microwave Transitions

Without a dedicated laser to drive transitions from the $|F = 1\rangle$ level, it was necessary to implement a scheme to use light resonant with the $|F = 2\rangle$ level to remove background atoms. Therefore, we included a system for driving microwave frequency transitions from $|1, 0\rangle$ to $|2, 0\rangle$.

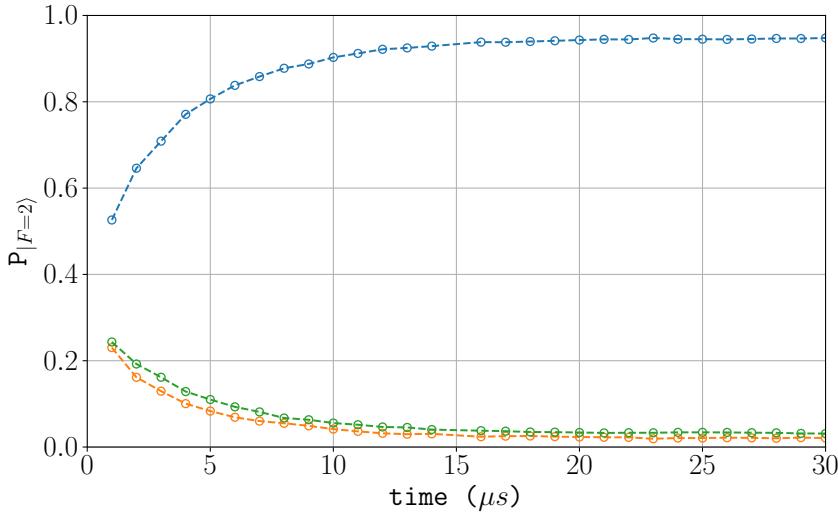


Figure 5.10: Population in each Zeeman sub-level as the $|F = 1\rangle \rightarrow |F' = 0\rangle$ pumping time is increased. The $m_F = 0, -1, +1$ populations are shown in blue, orange and green, respectively. After 30 μs , approximately 5% of the population remains in the $m_F = \pm 1$ sub-levels.

Microwave Generation

A diagram of the setup for this is shown in Figure 5.11. The microwave radiation is generated using a *Wind-Freak* synthesiser to output a microwave field oscillating at a frequency close to the hyperfine splitting frequency, $f_{\text{hfs}} = 6.838\,46\,\text{GHz}$. This is amplified by *MiniCircuits MCL ZRON-8G+* amplifier and directed into the chamber using a *Pasternack PE9859/SF-10* microwave horn, which produces a linearly-polarised microwave field. The horn was aligned to the chamber at the position which maximised the population of atoms in the $|2, 0\rangle$ state. The synthesiser is clocked using a stable 100 MHz signal from the μQuans laser. When the synthesiser was clocked using its internal 27 MHz reference clock, this produced a noticeable jitter in the output frequency, which led to a significant shot-to-shot fluctuation in the $|2, 0\rangle$ population.

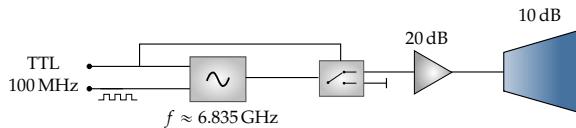


Figure 5.11: Schematic diagram of the microwave assembly. The frequency close to the hyperfine splitting frequency is generated by a *Wind-Freak* synthesiser. A 100 MHz clock signal acts as a stable reference frequency for the synthesiser. The generated microwave power is amplified twice, first by a low-power *Mini-Circuits* amplifier, then by a microwave horn, which produces a highly directional, linearly polarised wave. The output is controlled by a digital signal, both at the synthesiser and at a bi-directional microwave switch. The second port of this is blocked with a $50\ \Omega$ terminator to prevent reflections.

Pulse Characterisation

Figure 5.12 shows a measurement of the population in the $|F = 2\rangle$ level for increasing durations duration of the applied microwave pulse. Rabi oscillations between the $|1, 0\rangle$ and $|2, 0\rangle$ states are clearly present. The loss of coherence between the states can be explained by an inhomogeneous driving field. Once inside the chamber, the microwaves reflect and scatter off the interior surfaces which results in a spatially-dependent Rabi frequency. This also leads to a depolarisation of the field, as σ^\pm transitions were also observed. Initially, around 85% of the population was driven into $|F = 2\rangle$ using a microwave pulse of $100\ \mu\text{s}$. After improving the alignment of the magnetic field during the microwave pulse, this fraction increased to 97% - the remaining 3% being distributed across the $m_F = \pm 1$ states.

Figure ?? shows a spectrum obtained by varying the frequency of the microwave pulse. This shows the presence of $\Delta m = \pm 1$ transitions from $|1, 0\rangle$, as well as the fact that the population in $|1, m_F = \pm 1\rangle$ decreases after the $|F = 1\rangle \rightarrow |F' = 0\rangle$ pumping step is applied. The linewidth of the microwave transition is much narrower than the Zeeman splitting, so only the clock transition is driven when a pulse with a frequency close to f_{hfs} is applied.

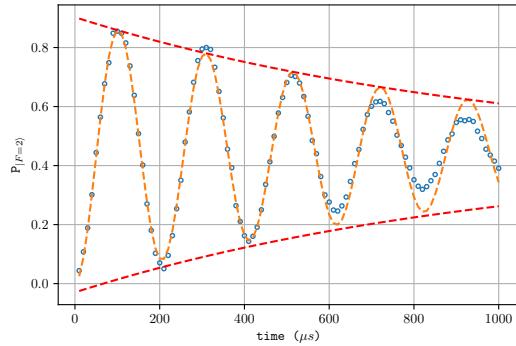


Figure 5.12: Damped Rabi oscillation between $|1,0\rangle$ and $|2,0\rangle$ using a microwave pulse of varying length. At longer pulse times, there is a loss of coherence due to a depahasing between the two states. The red dashed line is an envelope is a fit to a decaying exponential with a characteristic time of $\tau = 1016 \mu\text{s}$.

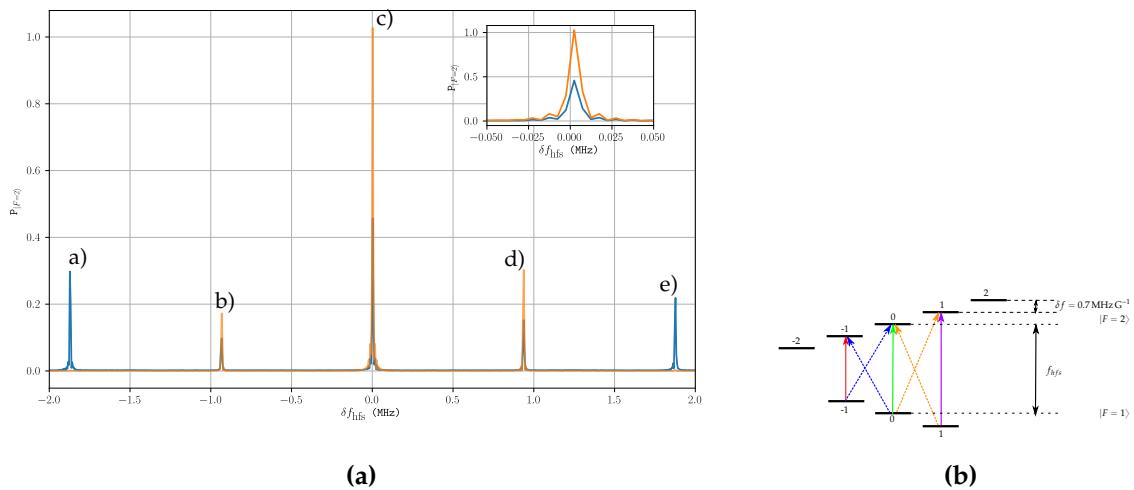


Figure 5.13: Microwave transition spectrum before (blue) and after (orange) $|F = 1\rangle \rightarrow |F' = 0\rangle$ pumping. (b) shows the transitions addressed as the microwave frequency is varied. Dashed and lines indicate $\Delta m = \pm 1$ transitions and solid lines indicate $\Delta m = 0$. In order of increasing frequency, the transitions in (a) are highlighted in: a) red, b) blue, c) green, d) orange and e) purple.

5.3.4 Blow-Away

After the atoms populate $|2, 0\rangle$, a velocity-selective Raman π -pulse is applied to transfer a fraction of those back into $|1, 0\rangle$. This step is discussed in detail in Section ???. For now, it suffices to say that a Raman pulse transfers 4% **Check this!** the atoms back to $|1, 0\rangle$. The remaining need to be removed, otherwise they contribute to a large background signal.

The final pulse during the state preparation sequence is used to push these non-contributing atoms out of the interferometer region. A single **MOT** beam is used so that there is a net momentum transfer to the atoms as they absorb light and fluoresce. The frequency of this blow-away beam is detuned from the $|F = 2\rangle \rightarrow |F' = 3\rangle$ transition by -3 MHz , which is the same frequency used for detection (see Section 6.4). A pulse of $50\text{ }\mu\text{s}$ is enough to remove all atoms in $|F = 2\rangle$.

5.3.5 Conclusion

This chapter has presented the stages of the experiment which are used to prepare an ensemble of atoms for interferometry. This requires cooling the atoms to limit the thermal expansion of the cloud during interferometry. The atoms are also launched using a moving molasses so that only one pair of beams is resonant with the Raman transition. Finally, we then apply a sequence of optical and microwave pulses, to increase the population of atoms in $|1, 0\rangle$. A velocity-selective Raman pulse with a narrow linewidth is used to make the velocity spread along the Raman axis much smaller than the Doppler width. Aside from some residual population in $|1, \pm 1\rangle$, the remaining atoms are removed using a pulse of light close to resonance with the $|F = 2\rangle \rightarrow |F' = 3\rangle$ transition. This results in an ensemble of which around 40% of the population contributes to the interferometer signal.

Chapter 6

Acceleration-Sensitive Interference

This chapter describes the work towards realising an atom interferometer and subsequently measuring accelerations.

6.1 Chapter Outline

To-Do:

- Raman spectrum, identifying each transition
- Characterisation of velocity-selective pulse and each interferometer pulse using Rabi oscillations.
- Making a three-pulse atom interferometer
- Improving acceleration sensitivity and correlating vibrations using MEMS

6.2 Raman Optical System

When designing an optical system for the light used in an atom interferometer, it is worth paying attention to both the spatial extent and beam waist of the collimated beam. These requirements are particularly important in this experiment, where acceleration due to gravity is perpendicular to the Raman beam axis and causes significant transverse motion of the atoms. Firstly, the optical system must be designed to make sure that the atoms are illuminated by each interferometer pulse. In addition to this, a more subtle requirement on the fringe contrast constrains the beam waist size. The gradient of intensity across the atom cloud must be small so that each atom is driven by (approximately) the same Rabi frequency. Otherwise, this variation in the Rabi frequency will dephase the atoms, which reduces the interferometer fringe contrast.

A further constraint on the optical system comes from the effect of thickness variation of optical elements. If the optical path length of the light as it passes through an element is not uniform, this will lead to wavefront aberrations. A spatially-varying phase leads to a bias in the interferometer phase since it does not depend on acceleration. Moreover, since this phase is not the same for each atom, it is another source of dephasing. Considering the effects of wavefront aberrations was a large motivating factor for designing an optical system for use inside the vacuum chamber. Specifically, the distortions of a laser wavefront through optical viewports would drastically reduce the fringe contrast under transverse motion of the atoms during the interferometer. The process used to bond an optical viewport to a flange stresses the glass and distorts its thickness, producing wavefront aberrations that factor into the phase uncertainty of an atom interferometer [Schkolnik2015]. **Add requirements on other optics**

6.2.1 Fringe Contrast Dependence

The effects of a gradient of intensity on the fringe contrast can be shown by considering an ensemble of atoms that are spatially distributed by a Gaussian distribution. Neglecting the effect of the ensemble's velocity distribution on the Raman detuning and for fixed pulse times, the pulse area $\Omega\tau$ varies only as a function of the radial displacement from the optic axis. The total fringe contrast can be determined by a convolution of the contrast for a single atom with the atomic density

$$\mathcal{C} = \int \frac{1}{\sqrt{2\pi}\sigma_c} e^{-r^2/(2\sigma_c^2)} f_{\pi/2-\pi-\pi/2}(\Omega(r-r_1), \Omega(r-r_2), \Omega(r-r_3)) dr \quad (6.1)$$

where σ_c is the radial width of the atom cloud, $f_{\pi/2-\pi-\pi/2}$ is the fringe contrast as previously described in equation (??) and r_i is the position of the ensemble's centre-of-mass at the i -th pulse. If the atom cloud is initially at the centre of the laser and falling under gravity, then these coordinates are $(0, -\frac{1}{2}gT^2, -2gT^2)$ respectively. Under the assumption that the two lasers which drive the Raman transition have the same waist size, the Rabi frequency, which is determined by the product of the electric fields (see equation (??)), can be described by

$$\Omega(r) = \Omega_0 e^{-2r^2/w^2} \quad (6.2)$$

where Ω_0 is the Rabi frequency along the optic axis and w is the waist size – the distance at which the electric field falls to $1/e$ of its peak value. The fringe contrast as a function as beam waist for an atom cloud of width a width $\sigma_c = 5$ mm and a time between interferometer pulses of $T = 25$ ms is plotted in Figure 6.1. For small beam waists, the intensity gradient across the cloud significantly reduces the fringe contrast. In fact, a beam waist much greater than the width of the cloud is necessary to achieve a large contrast between the two interferometer states. Relaxing the assumptions

made on the ensemble's velocity distribution to include its influence on the detuning and spatial distribution of the atoms during the interferometer would strengthen this argument.

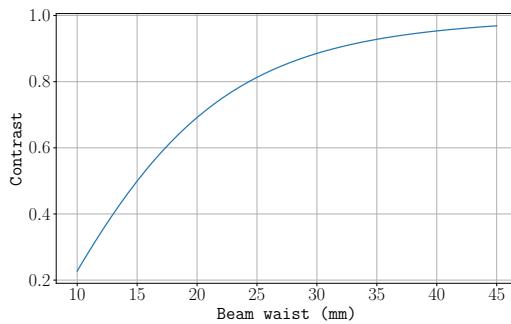


Figure 6.1: Simulated fringe contrast as a function of waist size w for an atom cloud falling under gravity. This model assumes a Gaussian distributed atomic density with a width $\sigma_c = 5$ mm and a time between interferometer pulses of $T = 25$ ms. For smaller beam waists the subsequent interferometer pulses have a larger intensity gradient across the atom ensemble, which increases the dephasing of the two states and reduces the interferometer fringe contrast.

So far, it has been shown that a large beam waist is necessary to achieve a high fringe contrast when allowing for transverse motion of the atoms across the laser wavefront. Otherwise, if the fringe contrast was poor, this would limit the sensitivity of the interferometer to accelerations rather than other effects which are less rectifiable. Another optical effect which influences the sensitivity is distortions of the laser wavefront. In an ideal case, the superposition of the spherical wavefronts of the two lasers results in a planar wavefront for the effective field which drives the Raman transition. However, propagation through rough optical elements distort these wavefronts and introduce a spatially varying component of the Raman phase that is independent of acceleration. If the atom cloud's trajectory is parallel with the Raman axis, then this additional phase is the same at each laser pulse and is therefore cancelled out. Of course, this does not occur when the cloud moves transverse to the Raman axis where this random phase has the effect of reducing the fringe contrast. Starting with the

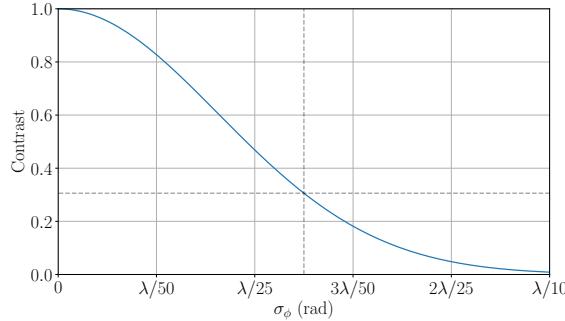


Figure 6.2: Expected contrast as a function of random phase contributions. This assumes that the phase imprinted on an atom during each interferometer pulse has an additional random component that is Gaussian distributed around 0 with a standard deviation of σ_ϕ . This random phase is also uncorrelated between each pulse so that the total can be obtained using Gaussian propagation of error. The dashed lines indicate the contrast for phase noise expected from conventional optics, which are usually engineered to a surface flatness of $\lambda/20$.

assumption that this phase is Gaussian distributed around 0, with a standard deviation of σ_ϕ , if this is uncorrelated at each interferometer pulse, then the interferometer phase $\Delta\Phi$ will be distributed with a standard deviation of $\sigma_\Phi = \sqrt{6}\sigma_\phi$. Denoting this random phase as $\delta\phi$, the fringe contrast is then given by

$$\mathcal{C}(\delta\phi) = \cos(2\delta\phi) \quad (6.3)$$

Following from this, if $\delta\phi$ is uncorrelated between each atom, the expected value of the contrast over the ensemble is given by

$$\langle \mathcal{C} \rangle = \frac{1}{\sqrt{2\pi}\sigma_\Phi} \int \mathcal{C}(\delta\phi) e^{-\delta\phi^2/2\sigma_\Phi^2} d\delta\phi \quad (6.4)$$

$$= e^{-2\sigma_\Phi^2} \quad (6.5)$$

The value of this expected contrast is plotted in Figure 6.2 and shows a strong dependence on σ_ϕ . Add more justification here. Cite Achim Peter's paper

6.2.2 Raman Beam Collimator

The optical system used to produce the beams for driving Raman transitions, which will conventionally be referred to as the Raman optics, was designed to reduce the previously mentioned effects which result in poorer interferometric fringe visibility and sensitivity to accelerations. Principally, the entire optical system was mounted inside the optical chamber so that the Raman light does not pass through any optical viewports before interacting with the atoms. Typically, the stress placed on the glass during the bonding process will distort the flatness more than is acceptable for achieving a high contrast. For example the viewports used for the MOT optics have a specified flatness of $\lambda/4$, so mounting the entire optical system inside the chamber was the simplest way to avoid a large distortion.

Figure 6.3a presents a diagram of the components used to send Raman light into the chamber and produce a collimated beam in the centre of the chamber. The light is coupled into the chamber using a UHV compatible PM fibre, manufactured by Diamond photonics. This is a kapton-coated PM-780 HP fibre that is bonded on one end to a DN16 flange using an epoxy resin. The external side of this flange has an FC/APC connector for coupling light from another fibre. Inside the chamber, the ferrule is connected to an FC/APC fibre plate. This is clamped between a piece which bolts onto the inside of a DN63 flange and another stainless steel plate which bolts onto the rest of the optics assembly. Fine adjustment of the position of the fibre along the optic axis is achieved using shim plates with a thickness ranging from 200–300 μm . The fibre plate is free to rotate so that the orientation of the fibre with respect to a QWP at the output of the collimator. This QWP is manufactured by Light Machinery, and is described further in Section 6.2.3. When the fibre is correctly orientated (e.g. when the slow axis of the fibre is at 45deg to the slow axis of the waveplate), the two Raman light fields are orthogonally circularly polarised.

The original design for the optical system consisted of a triplet lens, as a system of three lenses is capable of correcting for the five types of Seidel aberrations that distort rays of monochromatic light. This was designed and manufactured by IC Optical Systems. Another specification for this lens system was that it had to produce a collimated beam with a waist size of around 35 mm so that the sensitivity of the interferometer was not limited by the effects of intensity gradients across the atoms. Unfortunately, the triplet was designed with an incorrect Numerical Aperture (**NA**). With a focal length of 123.4 mm and a diameter of 50 mm, the triplet lens has a **NA** of 0.194. However, the nominal **NA** for PM780-HP fibre used in the UHV compatible **PM** fibre is 0.12. Consequently, the light from this fibre did not fill the **NA** of the triplet lens and produced a beam with a waist of 13mm. **Plot to illustrate this.** To address this issue, a pair of aspheric lenses was included to increase the divergence angle of light from the fibre. These are manufactured by Thorlabs and have a focal length of 4.51 mm (352230-B) and 15.29 mm (352260-B), respectively, to give a magnification of 3.39.

Alignment and Collimation

As one of the main motivations for mounting the Raman optics inside the vacuum chamber was to reduce the effects of wavefront distortions, it is worth highlighting how inaccurate alignment of the optics can lead to aberrations. As discussed before (see Section [6.2.1](#)), distortions of the wavefront leads to a dephasing and loss of interferometer fringe visibility. Here, the same figures of merit as before are used to consider what misalignment is acceptable to ensure that the phase of the Raman wavefront deviates by less than $\lambda/100$ after a transverse distance of 12.5 mm.

Taking the fibre as a point source, misalignment can occur if it is displaced from the front focal point of the optical system longitudinally along or transversely to the optic

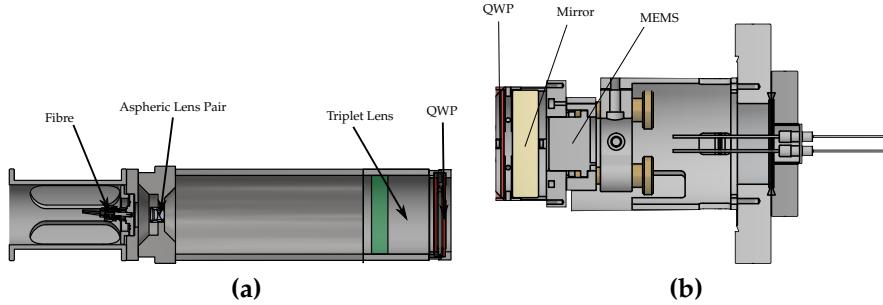


Figure 6.3: Diagrams of the components used in the Raman optical assemblies. (a) shows the collimator setup. Light is coupled into the chamber using a UHV fibre feedthrough. A pair of aspheric lenses is used to increase the divergence angle of the fibre output, before the light is collimated by a triplet lens. Finally, a quarter-wave plate is aligned so that it circularly polarises the collimated light fields. (b) illustrates the other half of the setup, which is used to retro-reflect the light. A second quarter-wave plate is used so that the reflected beams have the same handedness to their respective incoming ones. A MEMS accelerometer is mounted on the back of the mirror to measure vibrations. These components are all mounted on a piezo-controlled mirror mount whose tilt can be controlled from outside the vacuum chamber.

axis. If it is transversely displaced, this manifests as an angular displacement of the collimated light after the triplet lens. A large angular displacement is undesirable due to the fact that since one of the Raman light fields propagates further, the two wavefronts that drive acceleration-sensitive Raman transitions are not parallel. Figure 6.4a shows a simulation of the wavefront distortion as a result of this transverse misalignment. This is obtained by simulating the propagation of rays corresponding to each Raman light field through the optical system. The wavefront is estimated using the slope of each ray at a distance of 43 mm from the output of the triplet lens, which corresponds to the position of the centre of the vacuum chamber. The mirror is mounted at the same distance from the centre, so the second beam propagates 129 mm. Close to the optic axis, this distortion is approximately linear (i.e. a tilt) and it can be seen that a displacement of the fibre from the optic axis of <1 mm is sufficient to achieve the desired wavefront flatness.

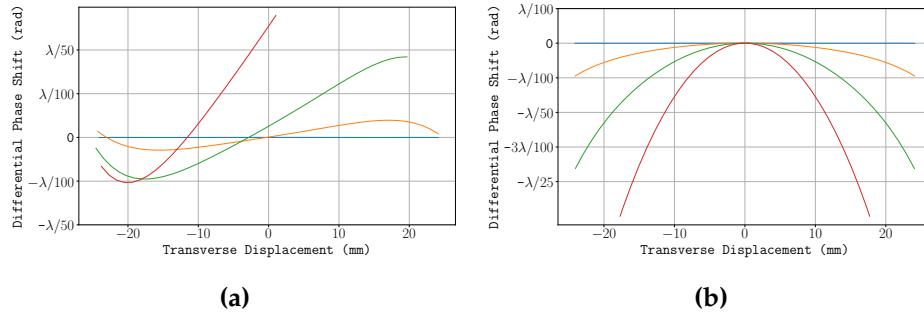


Figure 6.4: Simulated wavefront distortion for longitudinal and transverse fibre misalignment. Rays from a point source with a divergence angle corresponding to a NA of 0.12 are propagated through the Raman optical system. Rays corresponding to the reflected beam are propagated further with the assumption that the mirror is perpendicular to the optic axis. The first set of rays propagates 43 mm and the second propagates 129 mm. The wavefront for each beam is calculated by taking the slope of each ray and subtracting from the slope of the central ray. The wavefront of the effective field that drives the Raman transition is the difference of these two wavefronts. (a) shows the distortion of the wavefront for a transverse misalignment of the fibre for a displacement of 0 mm (blue), 0.5 mm (orange) 1 mm (green) and 1.5 mm (red) from the front focal point. (b) shows the wavefront for longitudinal displacements of 0 mm (blue), 0.3 mm (orange) 0.6 mm (green) and 1 mm (red).

Aside from a transverse displacement, it is possible that the fibre could be misaligned along the optic axis. In which case, the output beam will not be collimated. Consequently, the counter-propagating reflected rays will not be antiparallel to incoming ones. The effect of this longitudinal displacement on the Raman wavefront is shown in Figure 6.4b. Further from the optic axis the deviation in the phase of the light is greater, giving a quadratic distortion which is characteristic of a defocus. Comparing the wavefront distortion in this case, a requirement on the longitudinal misalignment of < 0.6 mm is needed for the previously specified flatness.

Measuring the Beam Width and Divergence

Most of the characterisation of the beam was taken outside of the vacuum chamber, prior to installation, due to extra difficulties of measuring its properties *in situ*. To measure the beam waist as well as its divergence, the beam was used to illuminate a

piece of paper aligned perpendicular to the optic axis. This was then imaged using a CCD camera with an objective lens so the entire spatial extent of the beam could be imaged. The camera was calibrated to give an effective pixel size of 5.1 pix/mm. Both the camera and paper were mounted on a translation stage, so that they could be moved along the direction of propagation. This

6.2.3 Retro-reflection Assembly

The mirror, also manufactured by Light Machinery, that is used to retro-reflect the incoming beams is mounted on the opposite side of the chamber. In order to drive Raman transitions (see Section 6.3 for further detail) with counter-propagating light fields, a second QWP, made to the same specifications as the one in front of the collimation optics, is mounted in front of the mirror. If the forward-propagating fields are circularly polarised, the reflected ones will be polarised with the same handedness regardless of the orientation of this second QWP.

During the manufacturing process, the waveplates and mirror were polished to reduce irregularities in the thickness of each QWP and the surface of the mirror. Figure 6.5 shows the variation in the thickness of the waveplate in front of the triplet lens, measured by Light Machinery using a white light interferometer. This has a standard deviation of 4.62 nm and corresponds a standard deviation of the optical path length of $8.6 \times 10^{-3}\lambda$.

The QWP and mirror are fixed onto the front plate of a UHV compatible MDI-HS mirror mount, manufactured by Radiant dye. The horizontal and vertical tilt of the mirror can be adjusted using two thumbscrew actuators which cause the front plate to pivot around a ball bearing. This mount is designed for applications which require high stability, but of course the alignment will still drift over time. To avoid the

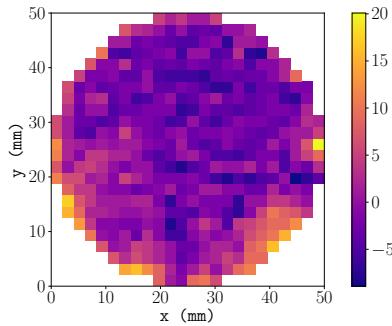


Figure 6.5: Thickness of the first **QWP!**, measured by a white light interferometer. The value is given in nm as a difference from the mean thickness. The standard deviation of this thickness is 4.62 nm and a peak-to-valley (PV) of need number here. Equivalent surface data for the other **QWP!** and mirror were not provided by Light Machinery, but had a PV thickness variation of 19 nm and 9 nm respectively.

need to periodically open the chamber to realign the mirror, a piezo-electric stack is placed between each actuator and the front plate so that the tilt of the mirror can be adjusted externally. Each piezo-stack is connected to a high-voltage feedthrough, so that their length (and hence mirror tilt) can be finely adjusted by controlling the voltage applied across them. A control voltage ranging between 0–10 V is amplified by a controller to give an applied voltage across the piezo stack that ranges between -10–150 V, corresponding to a travel range of 23 μm .

To understand the effect of misalignment, it is instructive to consider its effect on the effective wavevector \mathbf{k}_{eff} . As illustrated in Figure ??, if the mirror is misaligned from the incoming beam's wavevector by an angle θ , the two counter-propagating fields that drive Raman transitions have wavevectors $k_1(1, 0)$ and $k_2(\cos(\theta), \sin(\theta))$. $\mathbf{k}_{\text{eff}} = \mathbf{k}_1 - \mathbf{k}_2 \cos(2\theta)$. Fortunately, for small angular displacements, i.e. < 1 mrad, this does not greatly reduce the sensitivity to accelerations. In short, this means that \mathbf{k}_{eff} will have a spatially varying direction. Since an atom interacting via a Raman transition picks up a phase $\phi = \mathbf{k}_{\text{eff}} \cdot \mathbf{x}$, atoms travelling along different trajectories will accumulate different phases due to the spatial variation of \mathbf{k}_{eff} . Across the atom

ensemble, this leads to a dephasing and consequently, a loss of interferometer fringe visibility [10]

In-Situ Alignment and Optimisation

After mounting the Raman optical system inside the chamber, the mirror had to be aligned to retro-reflect the light. When the mirror is close to perpendicular to the light's wavevector, some of the power in the reflected beam couples back into the fibre. In principle, this power is maximised when the mirror is exactly perpendicular so maximising this power is a useful technique to align the mirror. A 99:1 fibre splitter was used to couple light into the chamber, which provided a means to measure the back-reflected power without needing any free-space optics. This was set up so that 99% of the incoming light entered the chamber, with the other 1% coupled into the corresponding output port. Due to the fact that a beam-splitter acts reversibly, 1% of the back-reflected light which couples into vacuum fibre exits the fibre-splitter on the other input port. Therefore, the power at this output was used to indirectly measure the alignment of the mirror.

Since the travel range of the piezo stacks does not cover the full motional range of the mirror mount, the mirror initially had to be manually aligned using the thumbscrew actuators. Once installed, the lack of direct access to optical system meant that conventional methods to coarsely align the mirror, such as observing the location of the reflected beam's focus, were not feasible. Rather than carry out the somewhat tedious job of systematically adjusting each thumbscrew until the mirror was aligned, an automatic routine was devised to do this. This was carried out using a pair of bipolar stepper motors that each rotated a ball driver inserted into the head of each thumbscrew. The revolution of these motors was controlled using an arduino microcontroller, which communicated to the computer using a serial interface. The motors

rotated by 0.9deg/step, which corresponds to a tilt of the mirror by 18.1 μ rad. This is smaller than the 0.67 mrad angular displacement that the piezo stack could provide, but the slow execution speed of the motor control meant that it was more practical to use a combination of the motors and piezos to systematically scan through the tilt of the mirror mount. **huge find out how big spot size was .**

Using this method, the mirror mount was aligned so that the maximum of the back-reflected power was reachable with the piezo stacks. Of course, it was foreseeable that the mirror would need to be periodically realigned, which would require another systematic iteration through the voltages applied to each piezo stack. Given that this search was quite time consuming, it was not a practical way to maintain alignment. To improve upon this, an optimisation method using the Nelder-Mead simplex algorithm [5] was implemented. This method is suitable for optimising multidimensional functions and has been used to demonstrate the automatic alignment of a fibre with up to 6 degrees of freedom [12]. In general terms, this algorithm aims to optimise the value of an objective function (in this instance, the optical power measured as a voltage by a photodiode) by sampling the function at various locations. For n parameters, a set of $n - 1$ points distributed randomly across the parameter space are chosen as the initial simplex. These are sorted in decreasing order of the value of the objective function and the algorithm proceeds by performing geometric transformations on this simplex, by sequentially reflecting, expanding and contracting this simplex. Each step starts with a reflection about the line between the two greatest values. The coordinates of the simplex are updated if the function has a greater value at the location given by one of these transformations, until the algorithm converges on a maximum value. As with many optimisation algorithms, the Nelder-Mead method has the potential to converge on a local optimum, but this is alleviated by expanding the simplex to look for more optimal values. The termination of the algorithm was decided by using the standard deviation of the last 5 values. Empirically, it was found that terminating

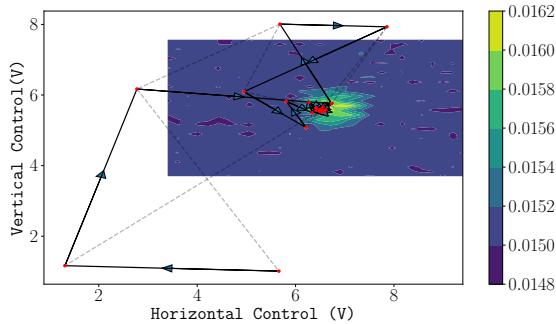


Figure 6.6: Automatic mirror alignment using the Nelder-Mead simplex algorithm. This procedure starts by randomly selecting three pairs of control voltages for the horizontal and vertical piezo stacks. At each co-ordinate, the back-reflected power is measured. The algorithm proceeds by geometrically transforming the simplex using reflections, expansions and contractions, and updating the simplex using this new co-ordinate if the power measured is greater than the current lowest value. The algorithm uses the standard deviation of the last 5 values as a check for convergence. In this case, it terminates once the standard deviation is smaller than $10 \mu\text{V}$. The shaded lines indicate the simplex bounded by the three co-ordinates at each iteration, whose area reduces as the algorithm converges on the optimum value. A scan of the piezo control voltages close to the optimum is also plotted. The irregular shape of the measured power is a result of a hysteresis effect when the horizontal control voltage was changed from its maximum value to the minimum. Even with a low signal-to-noise ratio, the algorithm converged on a value close to the optimum. This resulted in a misalignment of less than $13 \mu\text{rad}$ along both axes.

when the standard deviation was less than $10 \mu\text{V}$ resulted in stable performance of the algorithm, even when the signal-to-noise ratio of the measured voltage was poor. An example of this algorithm aligning the mirror mount is presented in Figure 6.6. To verify that the converged value was optimal, a systematic scan of the piezo stack control voltages in the region around this value was also carried out. In this case, the algorithm converged on a local maximum, but one that greatly enhanced the coupling efficiency of the reflected light back into the fibre. The difference in the piezo control voltages from their optimal values corresponds to a tilt of the mirror mount along the horizontal and vertical axis of less than $13 \mu\text{rad}$.

6.2.4 The MEMS Accelerometer

6.3 Driving Raman Transitions

As the forward-propagating fields have opposite handedness, the counter-propagating fields have the same handedness as the field

6.3.1 Frequency and Phase Control

6.4 Atom Detection

6.4.1 Optical System

6.4.2 Measuring the Interferometer Phase

6.5 Individual Pulse Characterisation

6.5.1 Velocity-Selective Pulse

6.5.2 Interferometer Pulses

6.6 Three-Pulse Atom Interference

6.6.1 Fringe Contrast vs. Launch Trajectory

6.7 Measuring Accelerations

6.7.1 Vibration Sensitivity

Chapter 7

Outlook

This final chapter describes some of the next steps and further work

7.1 Combining with classical accelerometers

- Discuss schemes for combining multiple sensors - Kalman filtering
- Extend this to inertial navigation
- Steps towards overcoming sensitivity-bandwidth trade-off.

7.2 Extending to sensitivity along three axes

- New chamber design
- Improvements to MSquared laser
- Required knowledge of gravitational axis for accurate navigation

Bibliography

- [1] D. J. Berkeland and M. G. Boshier. "Destabilization of dark states and optical spectroscopy in Zeeman-degenerate atomic systems". In: *Physical Review A* 65.3 (2002), p. 033413. ISSN: 1050-2947. DOI: [10.1103/PhysRevA.65.033413](https://doi.org/10.1103/PhysRevA.65.033413).
- [2] K. Dieckmann et al. "Two-dimensional magneto-optical trap as a source of slow atoms". In: *Physical Review A* 58.5 (1998), pp. 3891–3895. ISSN: 1050-2947. DOI: [10.1103/PhysRevA.58.3891](https://doi.org/10.1103/PhysRevA.58.3891).
- [3] Aviv Keshet and Wolfgang Ketterle. "A Distributed GUI-based Computer Control System for Atomic Physics Experiments". In: (2012). DOI: [10.1063/1.4773536](https://doi.org/10.1063/1.4773536). arXiv: [1208.2607](https://arxiv.org/abs/1208.2607).
- [4] *μQuans Laser System Specifications*.
- [5] J. A. Nelder and R. Mead. "A Simplex Method for Function Minimization". In: *The Computer Journal* 7.4 (1965), pp. 308–313. ISSN: 0010-4620. DOI: [10.1093/comjnl/7.4.308](https://doi.org/10.1093/comjnl/7.4.308).
- [6] Shin-ichi Ohshima et al. "Cesium Atomic Fountain with Two-Dimensional Moving Molasses". In: *Japanese Journal of Applied Physics* 34.Part 2, No. 9A (1995), pp. L1170–L1173. ISSN: 00214922. DOI: [10.1143/JJAP.34.L1170](https://doi.org/10.1143/JJAP.34.L1170).
- [7] M. Prentiss et al. "Atomic-density-dependent losses in an optical trap". In: *Optics Letters* 13.6 (1988), p. 452. ISSN: 0146-9592. DOI: [10.1364/OL.13.000452](https://doi.org/10.1364/OL.13.000452).

-
- [8] J. Schoser et al. "Intense source of cold Rb atoms from a pure two-dimensional magneto-optical trap". In: *Physical Review A* 66.2 (2002), p. 023410. ISSN: 1050-2947. DOI: [10.1103/PhysRevA.66.023410](https://doi.org/10.1103/PhysRevA.66.023410).
 - [9] A. M. AM Steane, M. Chowdhury, and CJ J. Foot. "Radiation force in the magneto-optical trap". In: *Journal of the Optical Society of America B* 9.12 (1992), p. 2142. ISSN: 0740-3224. DOI: [10.1364/JOSAB.9.002142](https://doi.org/10.1364/JOSAB.9.002142).
 - [10] G Tackmann et al. "Self-alignment of a compact large-area atomic Sagnac interferometer". In: *New Journal of Physics* 14.1 (2012), p. 015002. ISSN: 1367-2630. DOI: [10.1088/1367-2630/14/1/015002](https://doi.org/10.1088/1367-2630/14/1/015002).
 - [11] M. Walhout et al. " $\sigma_- + \sigma_\perp$ Optical molasses in a longitudinal magnetic field". In: *Journal of the Optical Society of America B* 9.11 (1992), p. 1997. ISSN: 0740-3224. DOI: [10.1364/JOSAB.9.001997](https://doi.org/10.1364/JOSAB.9.001997).
 - [12] R. Zhang and F.G. Shi. "A Novel Algorithm for Fiber-Optic Alignment Automation". In: *IEEE Transactions on Advanced Packaging* 27.1 (2004), pp. 173–178. ISSN: 1521-3323. DOI: [10.1109/TADVP.2004.825434](https://doi.org/10.1109/TADVP.2004.825434).

Acronyms

CCM Centre for Cold Matter

⁸⁷Rb Rubidium-87

⁸⁵Rb Rubidium-85

MOT Magneto-optical Trap

VSR Velocity-selective resonance

SDLC sub-Doppler laser cooling

AOM Acousto-optic Modulator

EOM Electro-optic Modulator

PM Polarisation-Maintaining

QWP Quarter-wave Plate

HWP Half-wave Plate

MFD Mode Field Diameter

NA Numerical Aperture

PPLN Periodically Poled Lithium Niobate

PLL Phase-Locked Loop

FPGA Field-Programmable Gate Array

EDFA Erbium-Doped Fibre Amplifier

ECDL External-Cavity Diode Laser

TTL Transistor-transistor Logic Circuit

NI National Instruments

DAQ Data Acquisition

VCO Voltage-Controlled Oscillator

ADC Analogue-to-Digital Converter

DAC Digital-to-Analogue Converter

HAL Hardware Abstraction Layer

SPI Serial Programming Interface

DDS Direct Digital Synthesiser

PBS Polarising beam-splitter

DRO Dielectric Resonator Oscillator

